



Microsoft Ignite The Tour

Learn. Explore. Connect.

London, England






An Introduction to SQL Server 2019

- Big Data Clusters


Paul Andrew

Data Platform MVP & Solution Architect




What

Why





How



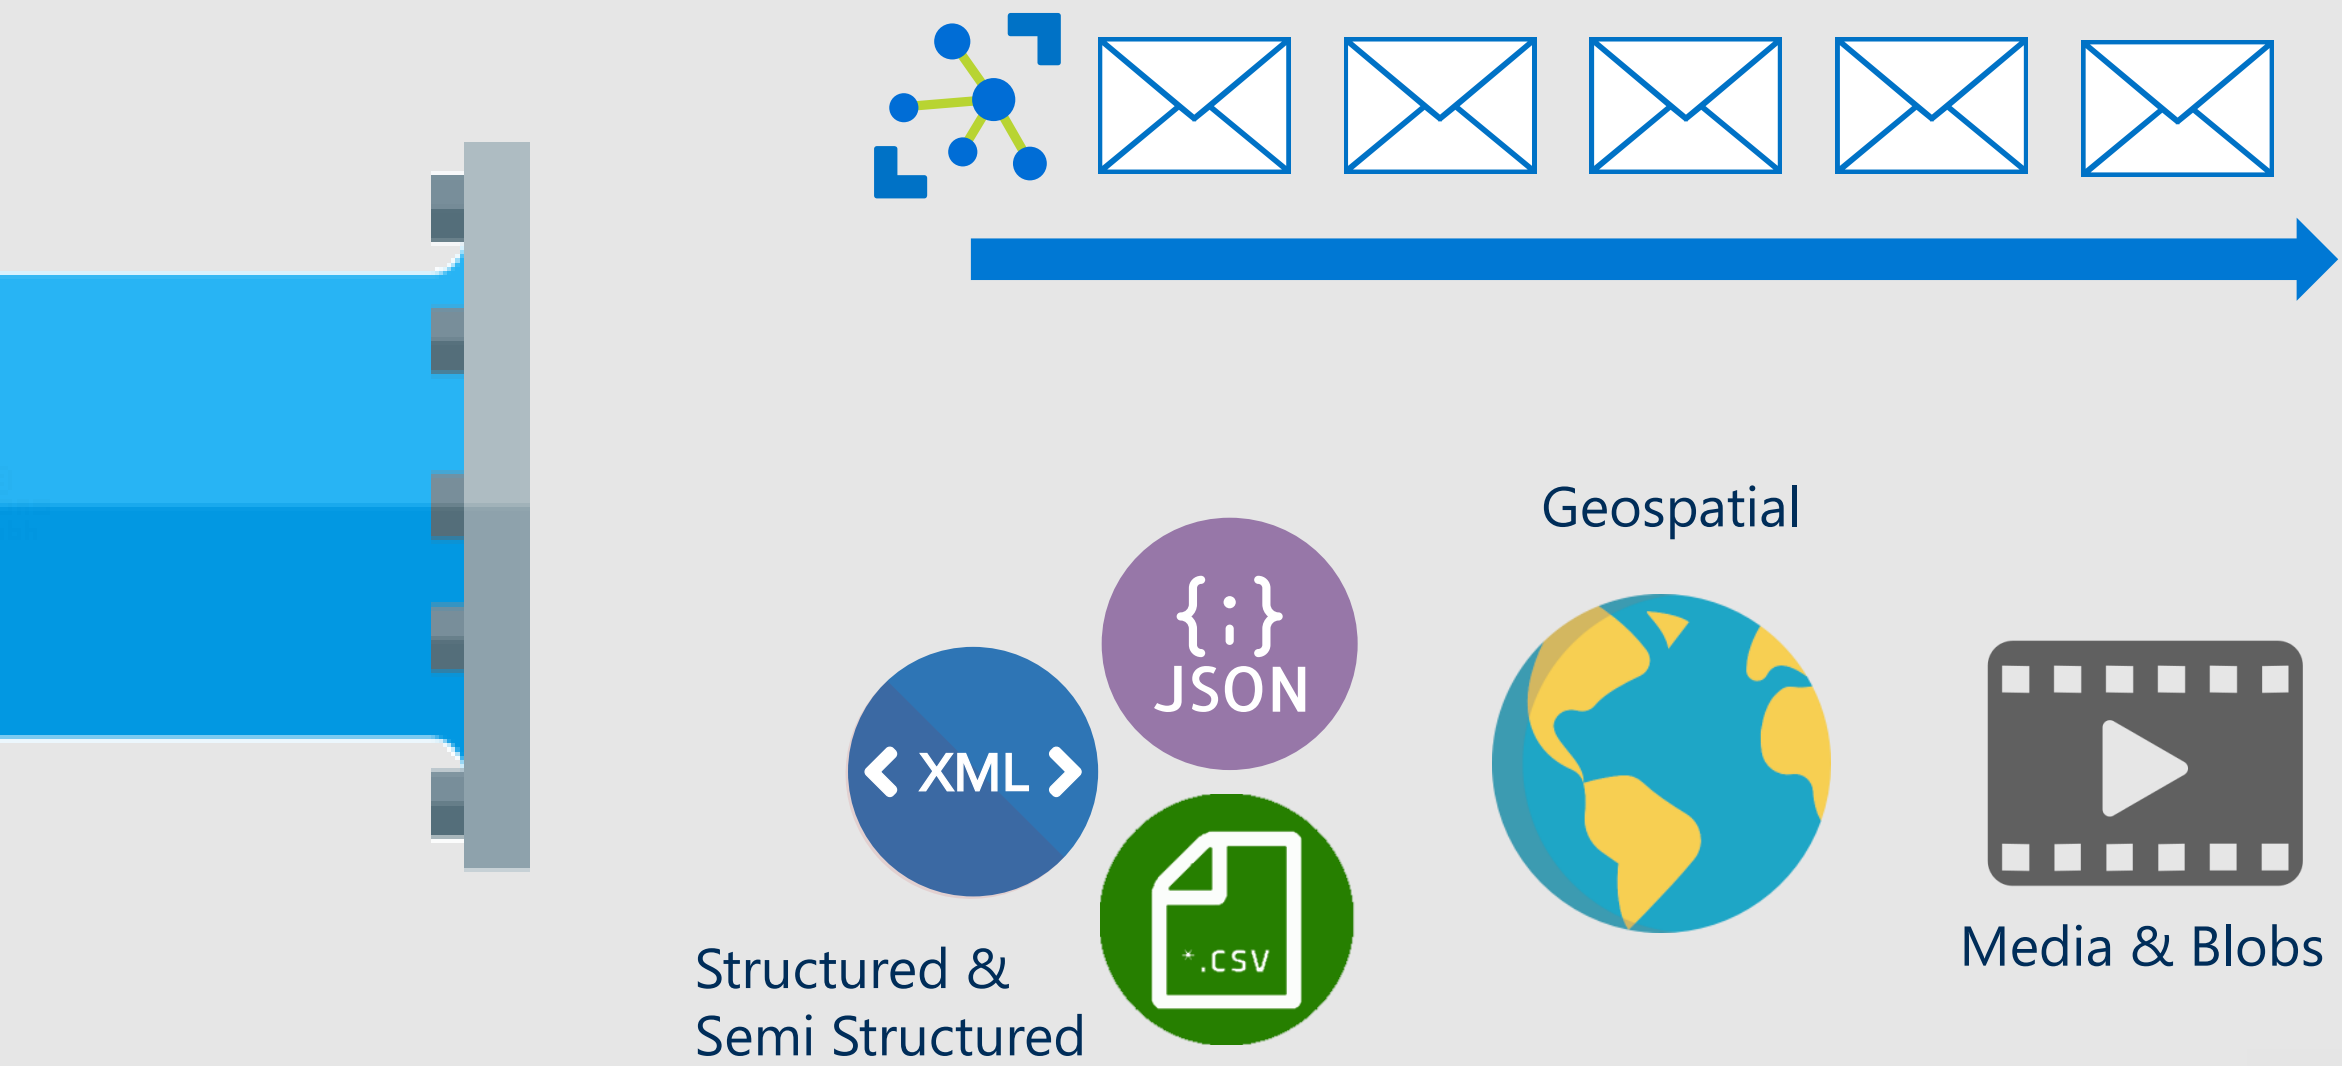
What

Why

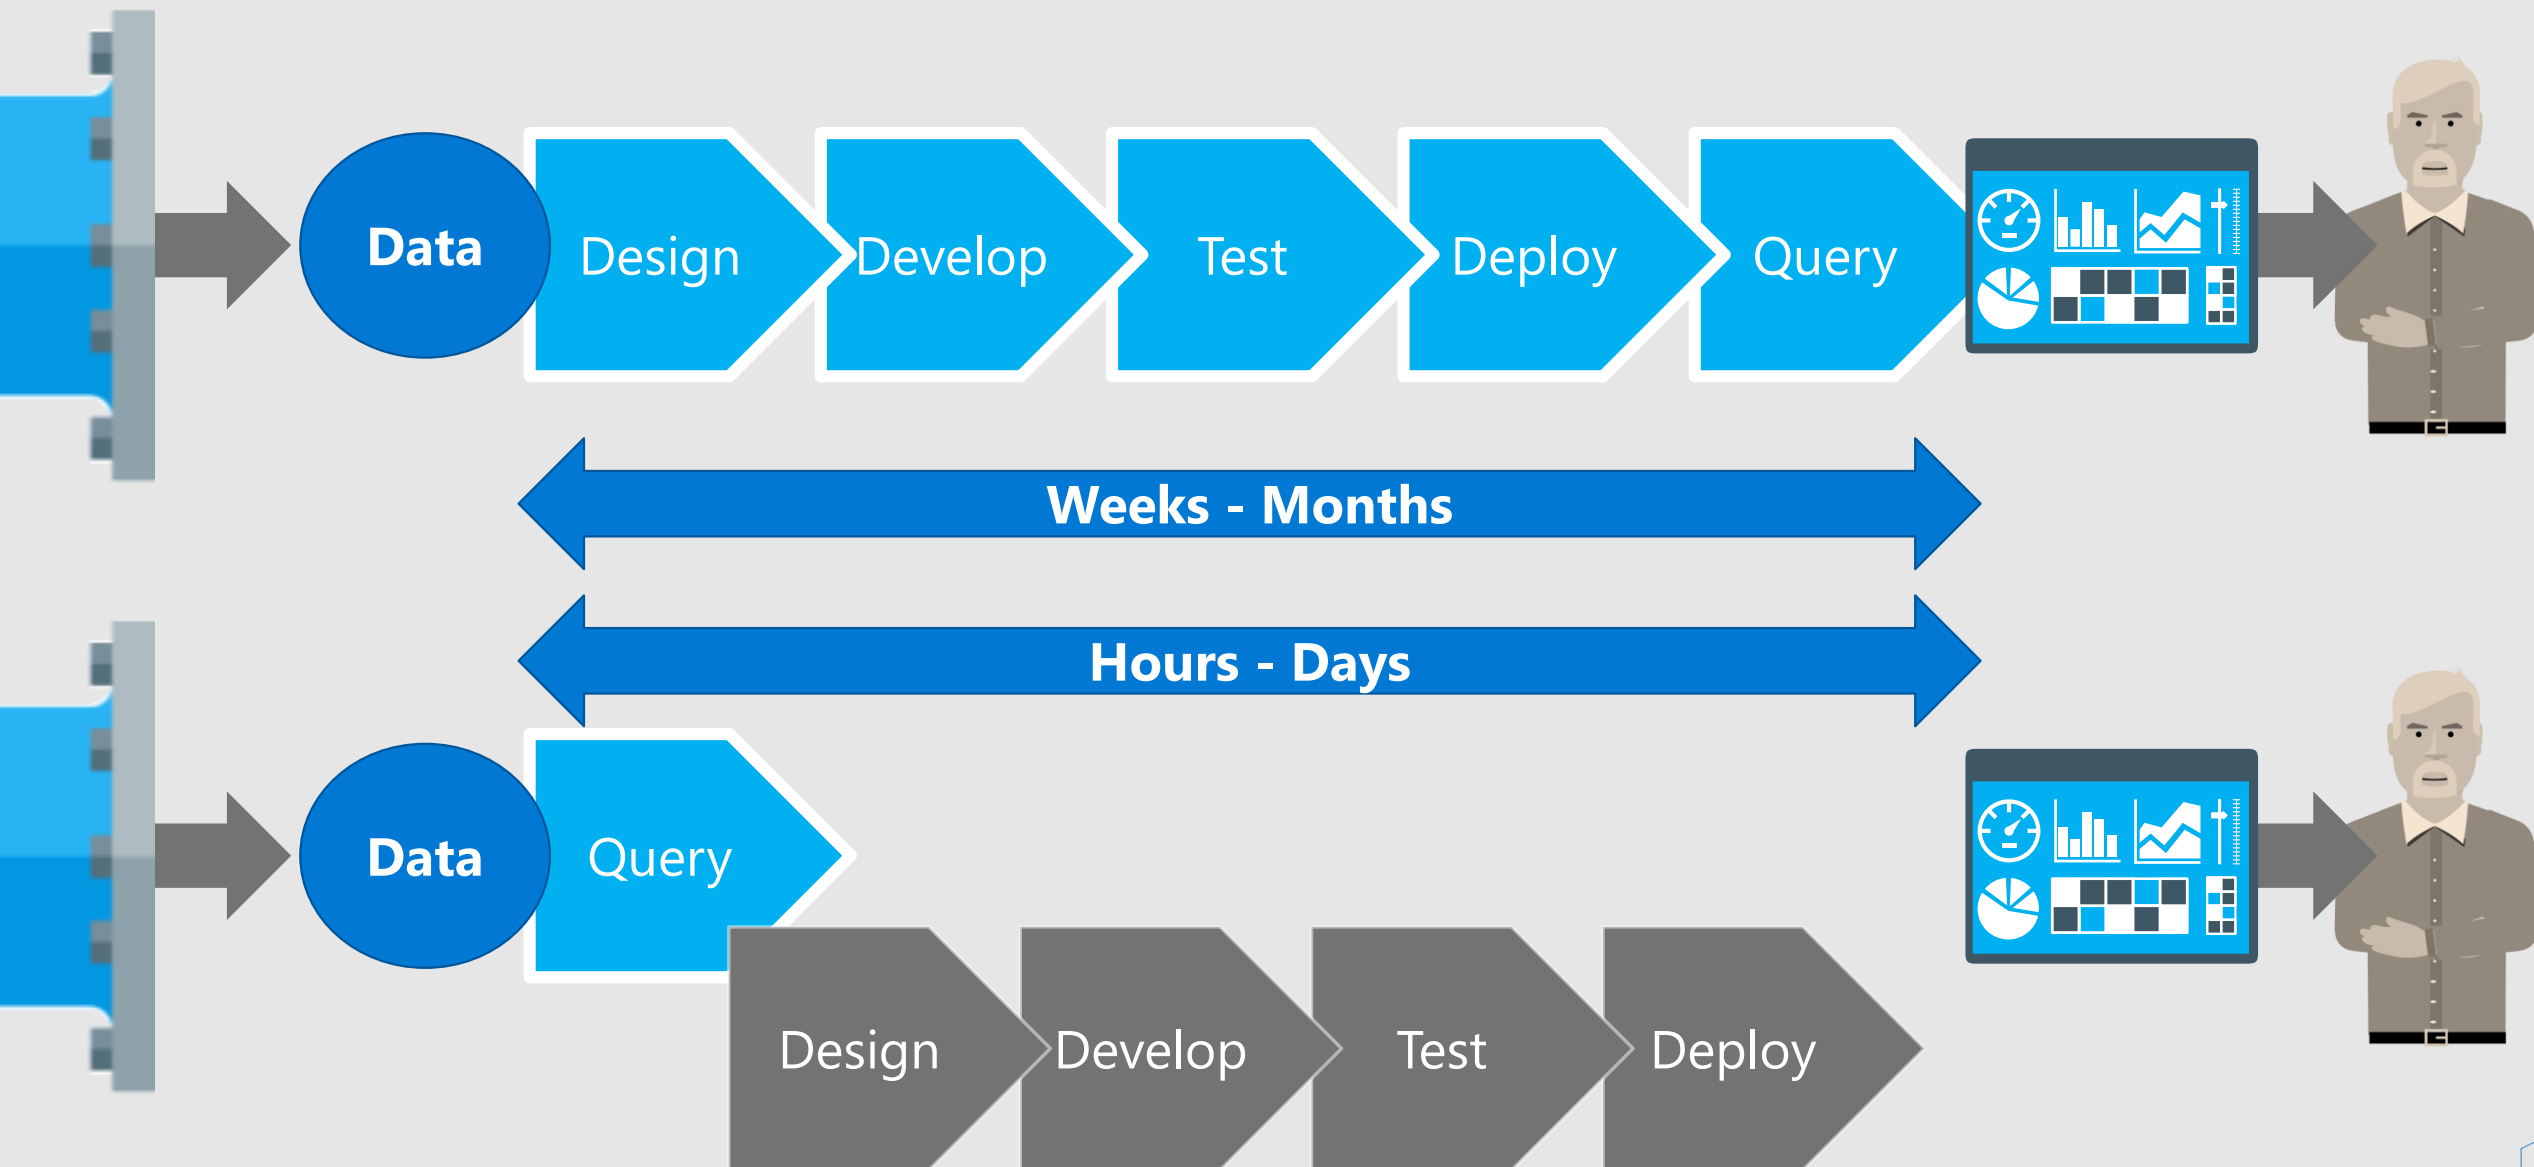
How



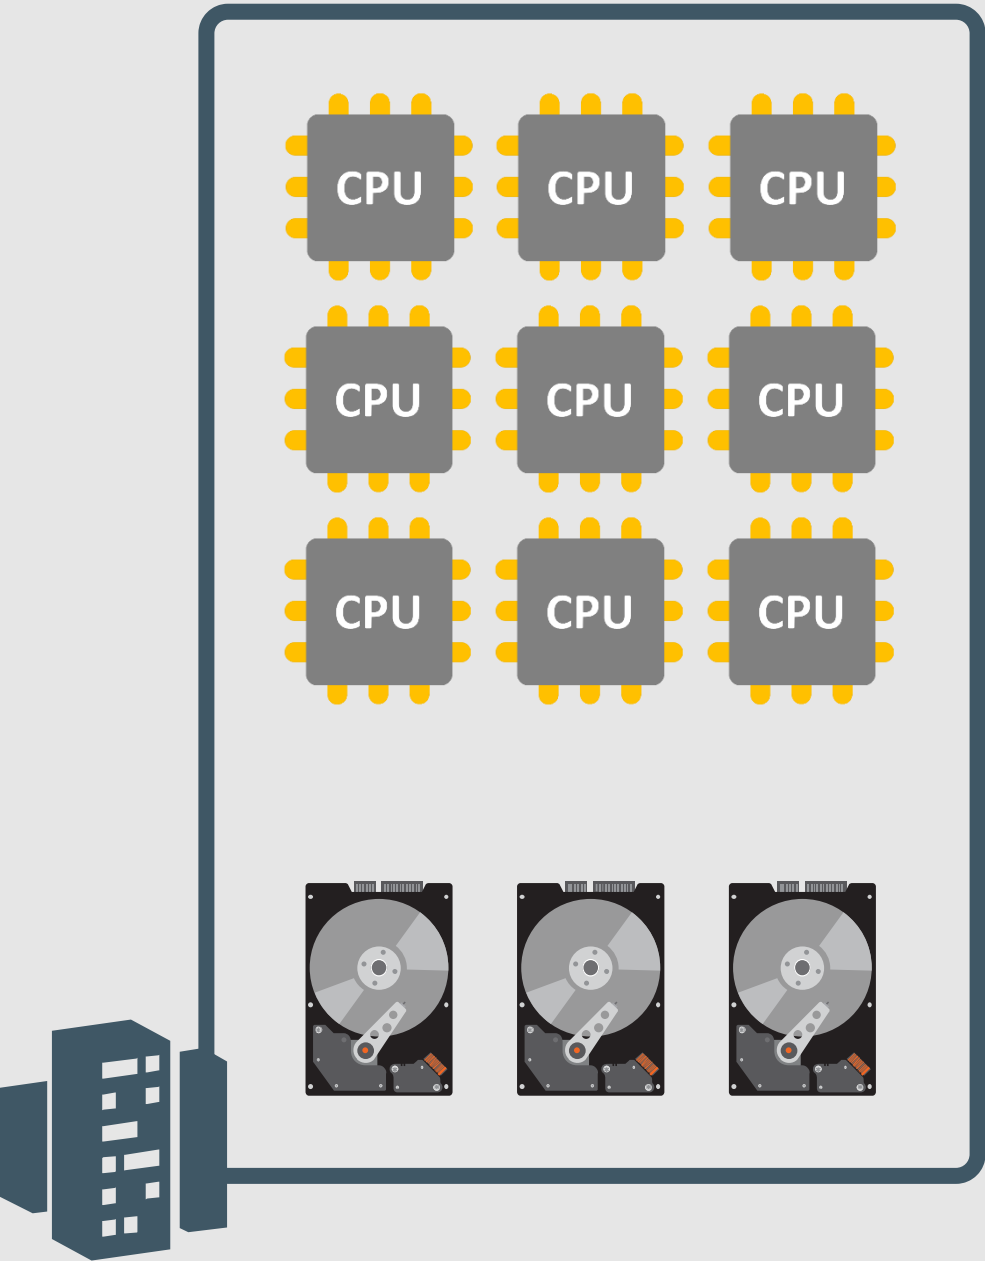
Problem 1: All Forms of Data Structure





Problem 2: Quick Actionable Insights



Problem 3: Infrastructure Costs & Scalability



Daily Server Utilization




The Big Data V's

Volume



Velocity

Variety


Veracity




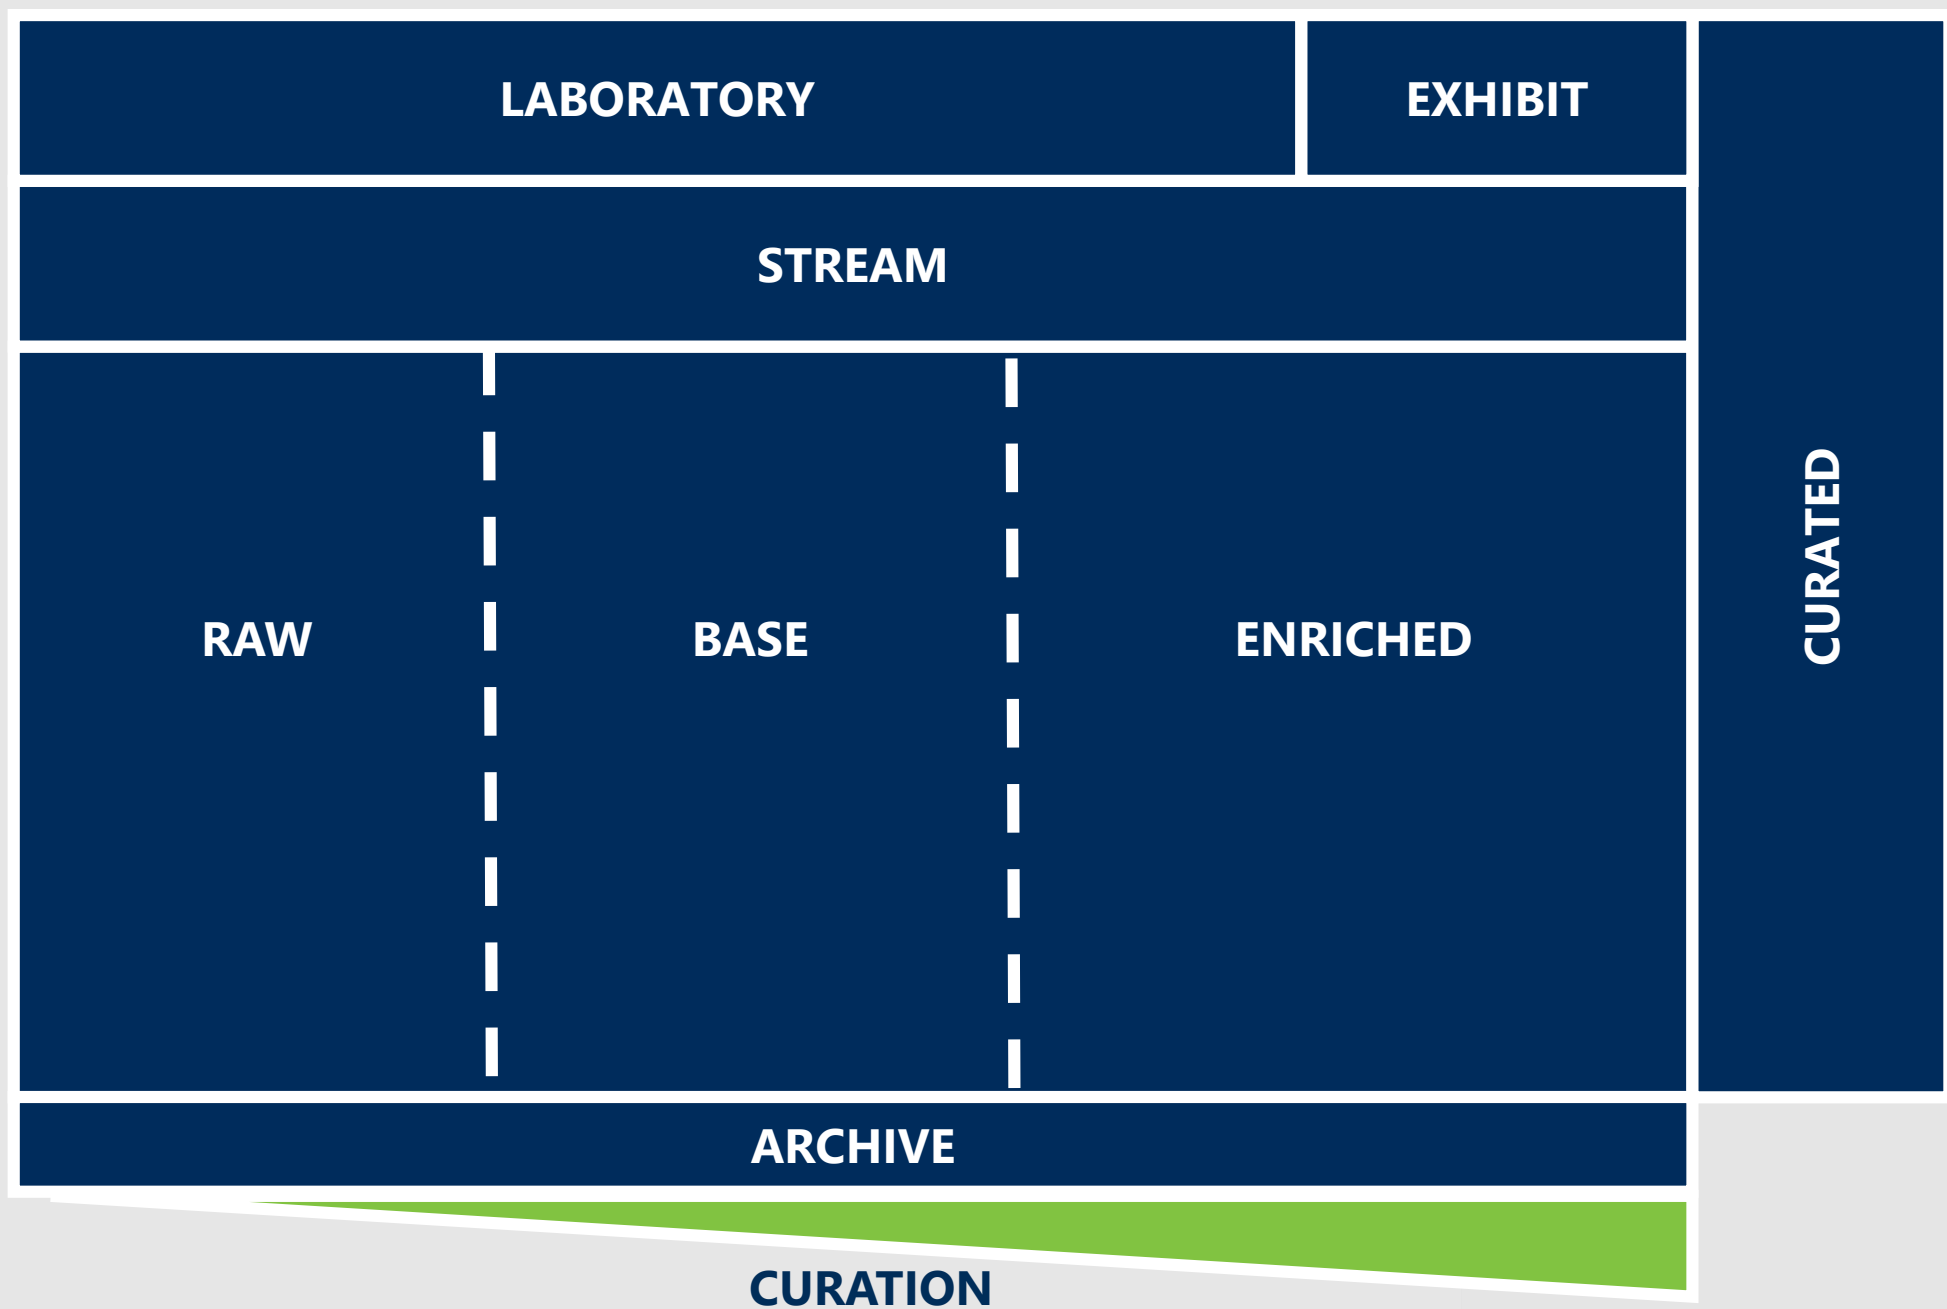
Data Lakes




Data Lakes





Data Lake – with Adatis Framework




The Modern Data Warehouse

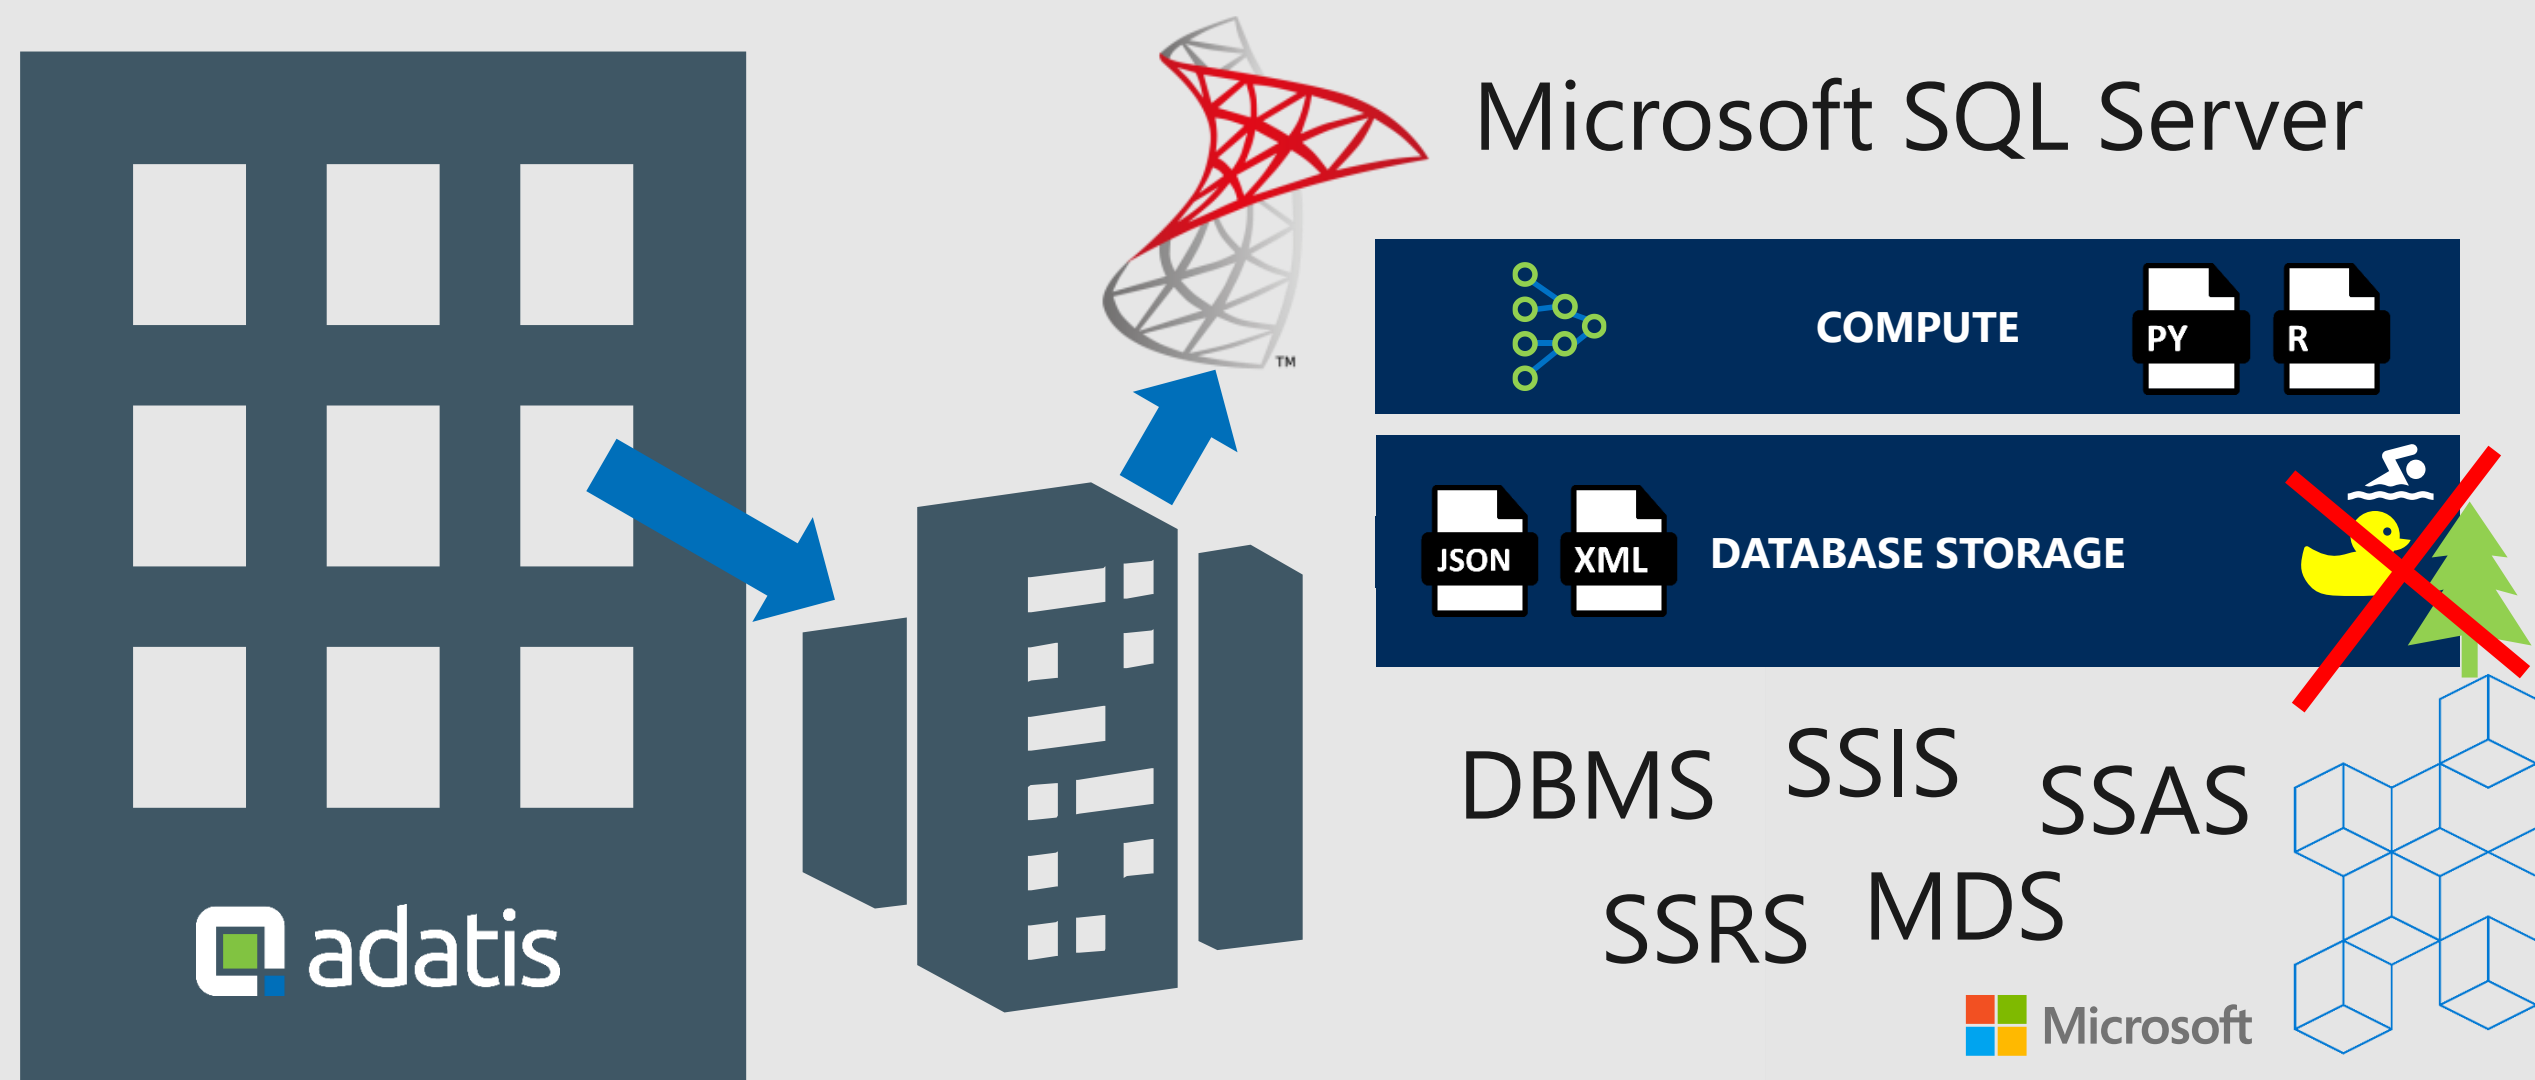


Cloud

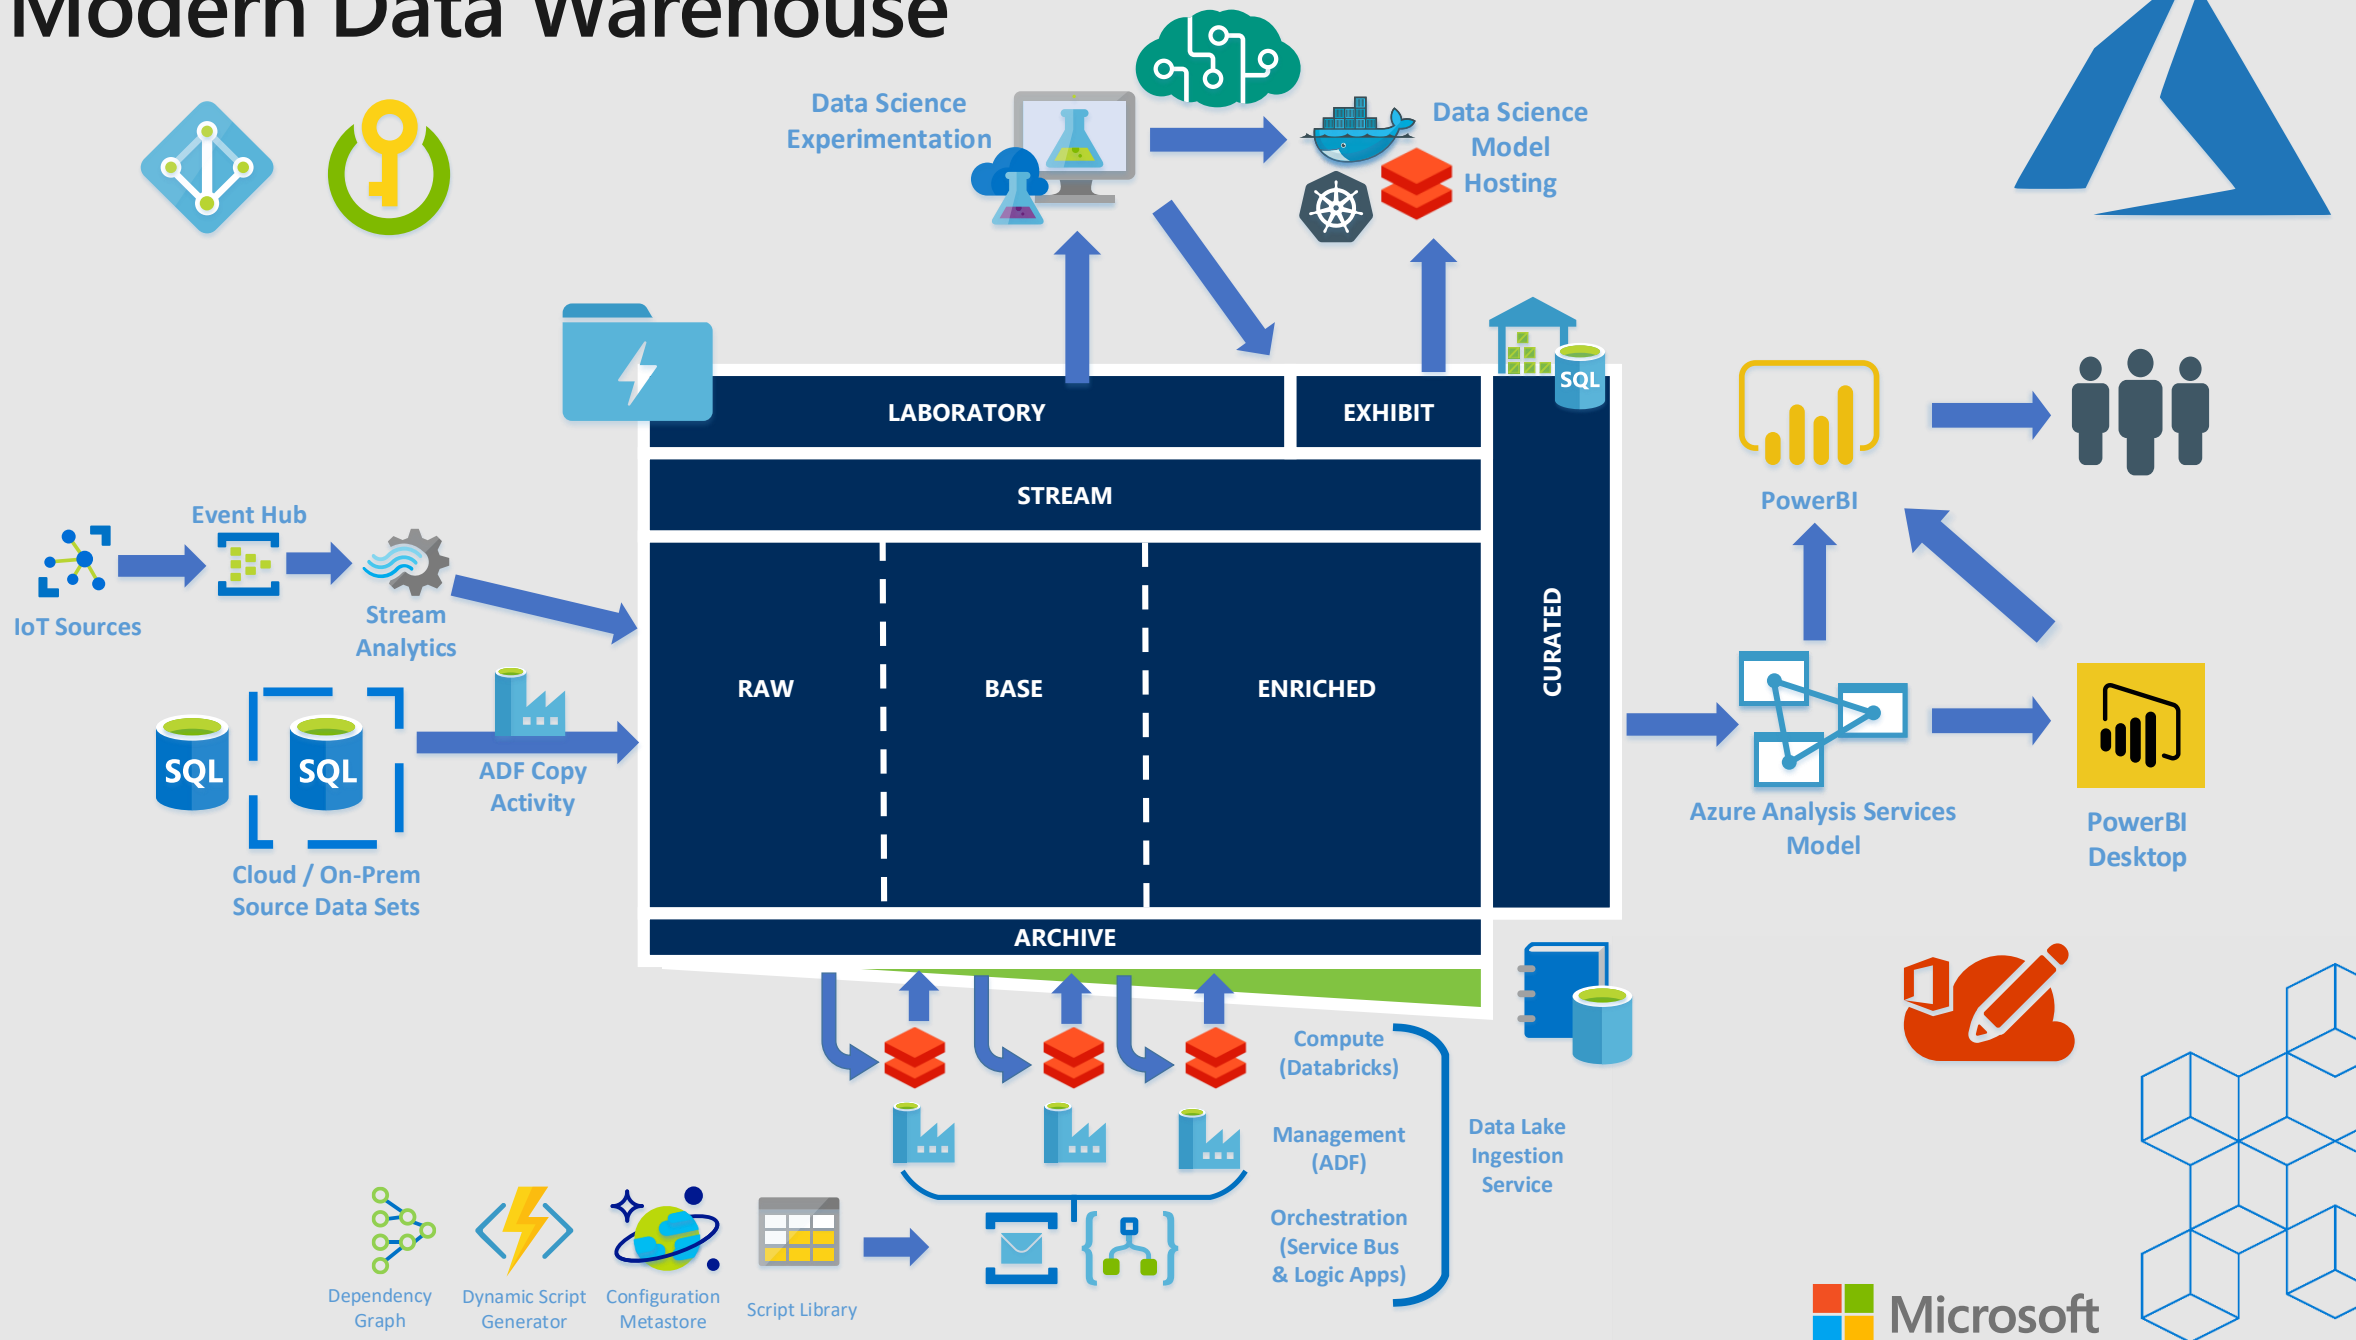




On Premises





The Modern Data Warehouse





What

Why

How




SQL Server - Product Evolution



Azure SQL Database (SQLDB)




Compute




Scale Up Compute

Azure SQL Data Warehouse (SQLDW)

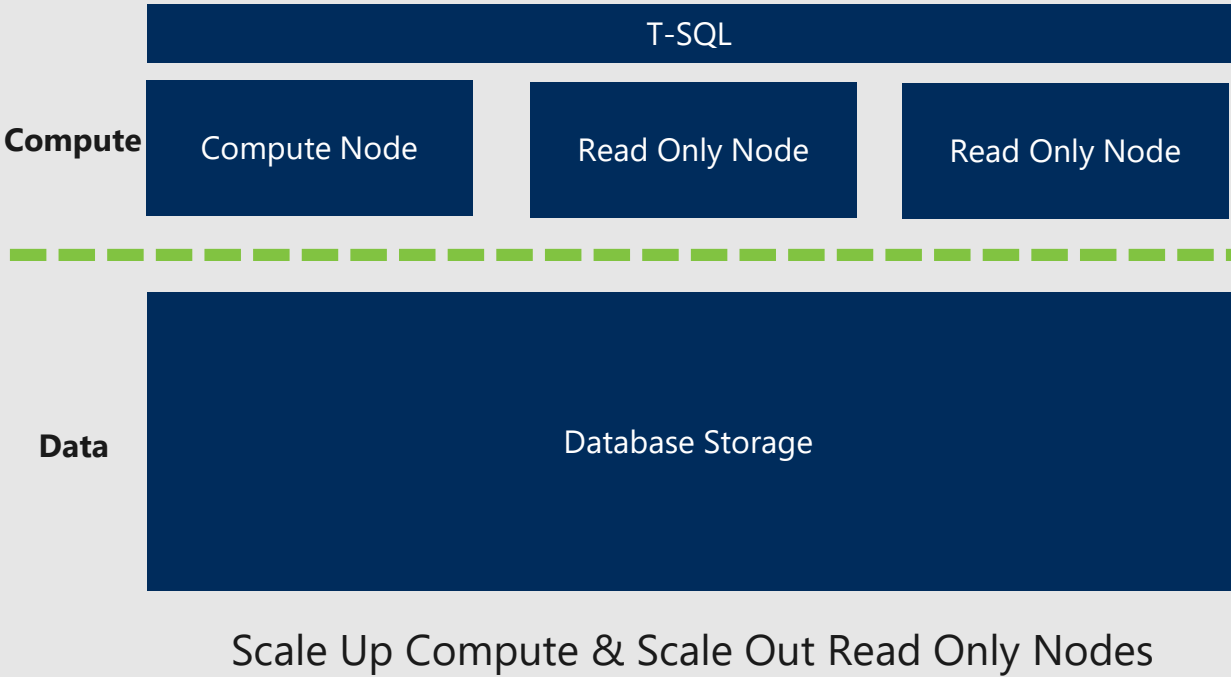


Scale Out Compute & Distributed Storage


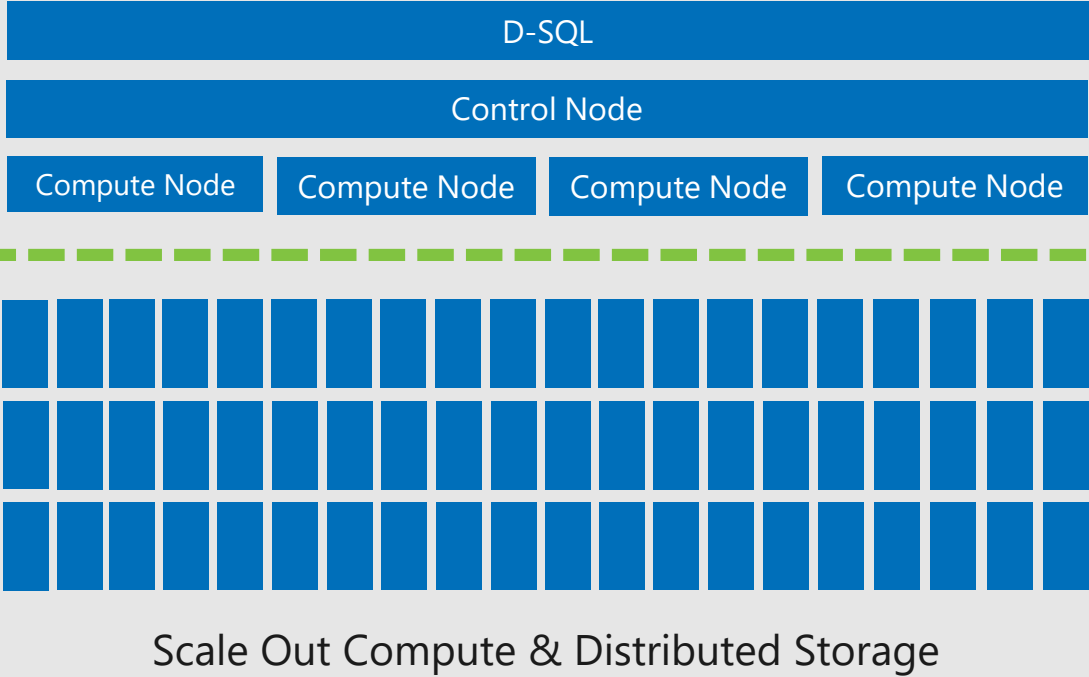
SQL Server - Product Evolution



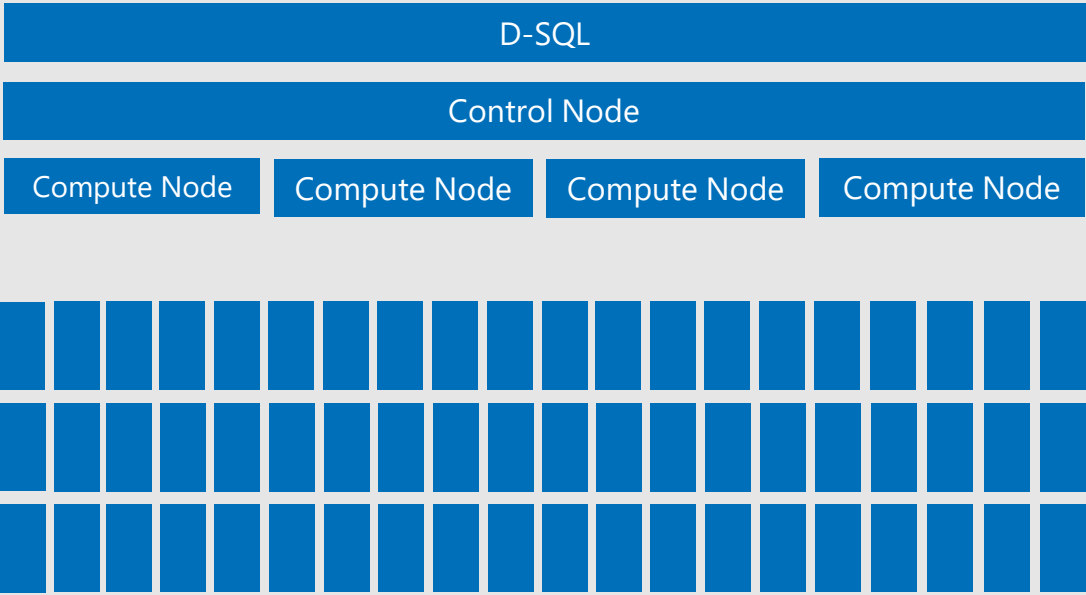
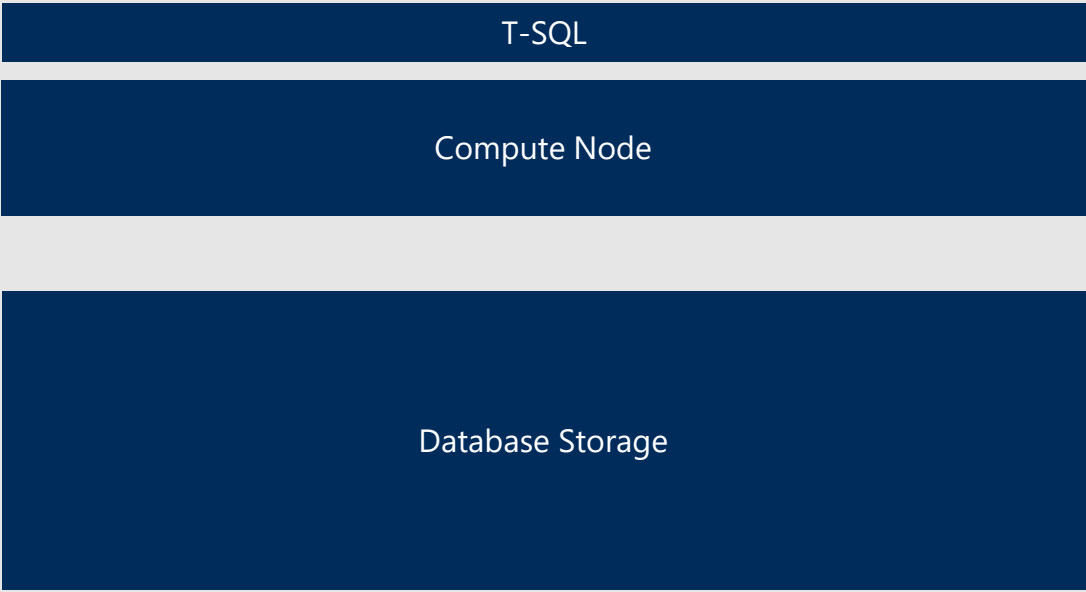

Azure SQL Database (SQLDB)




Azure SQL Data Warehouse (SQLDW)



SQL Server - Product Evolution






SQL Server - Product Evolution



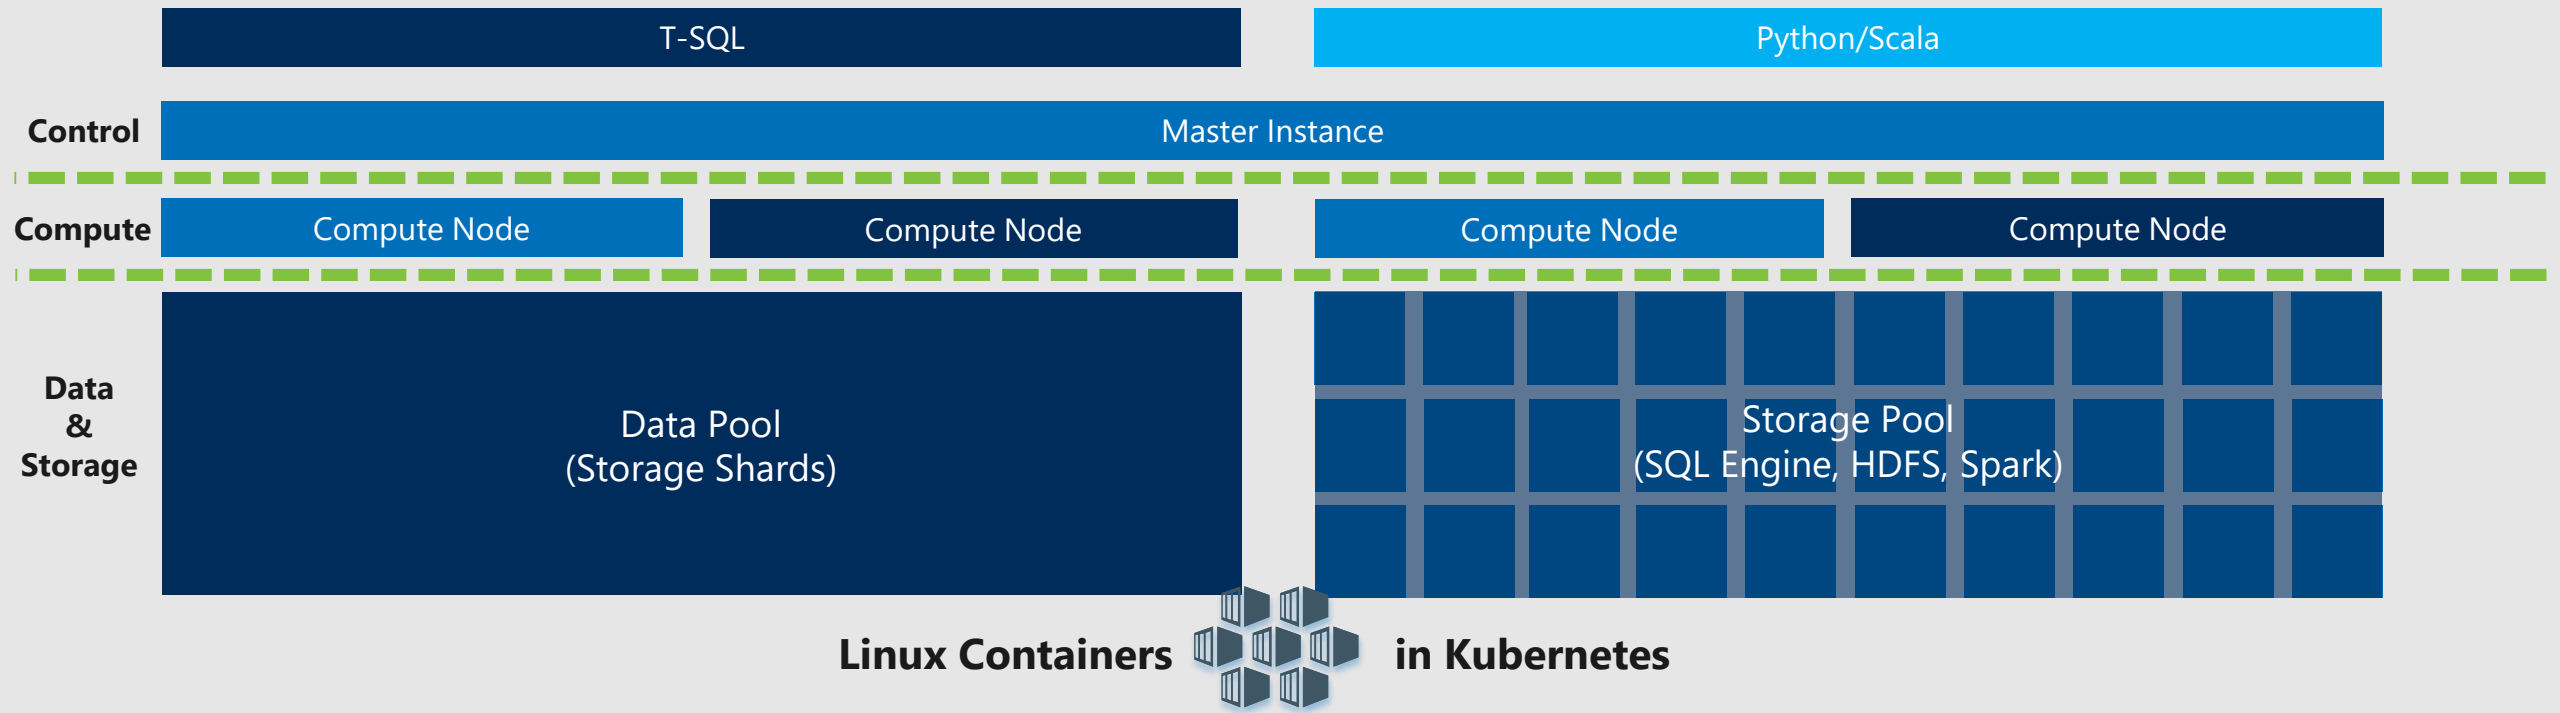
SQL Server 2019 Big Data Cluster



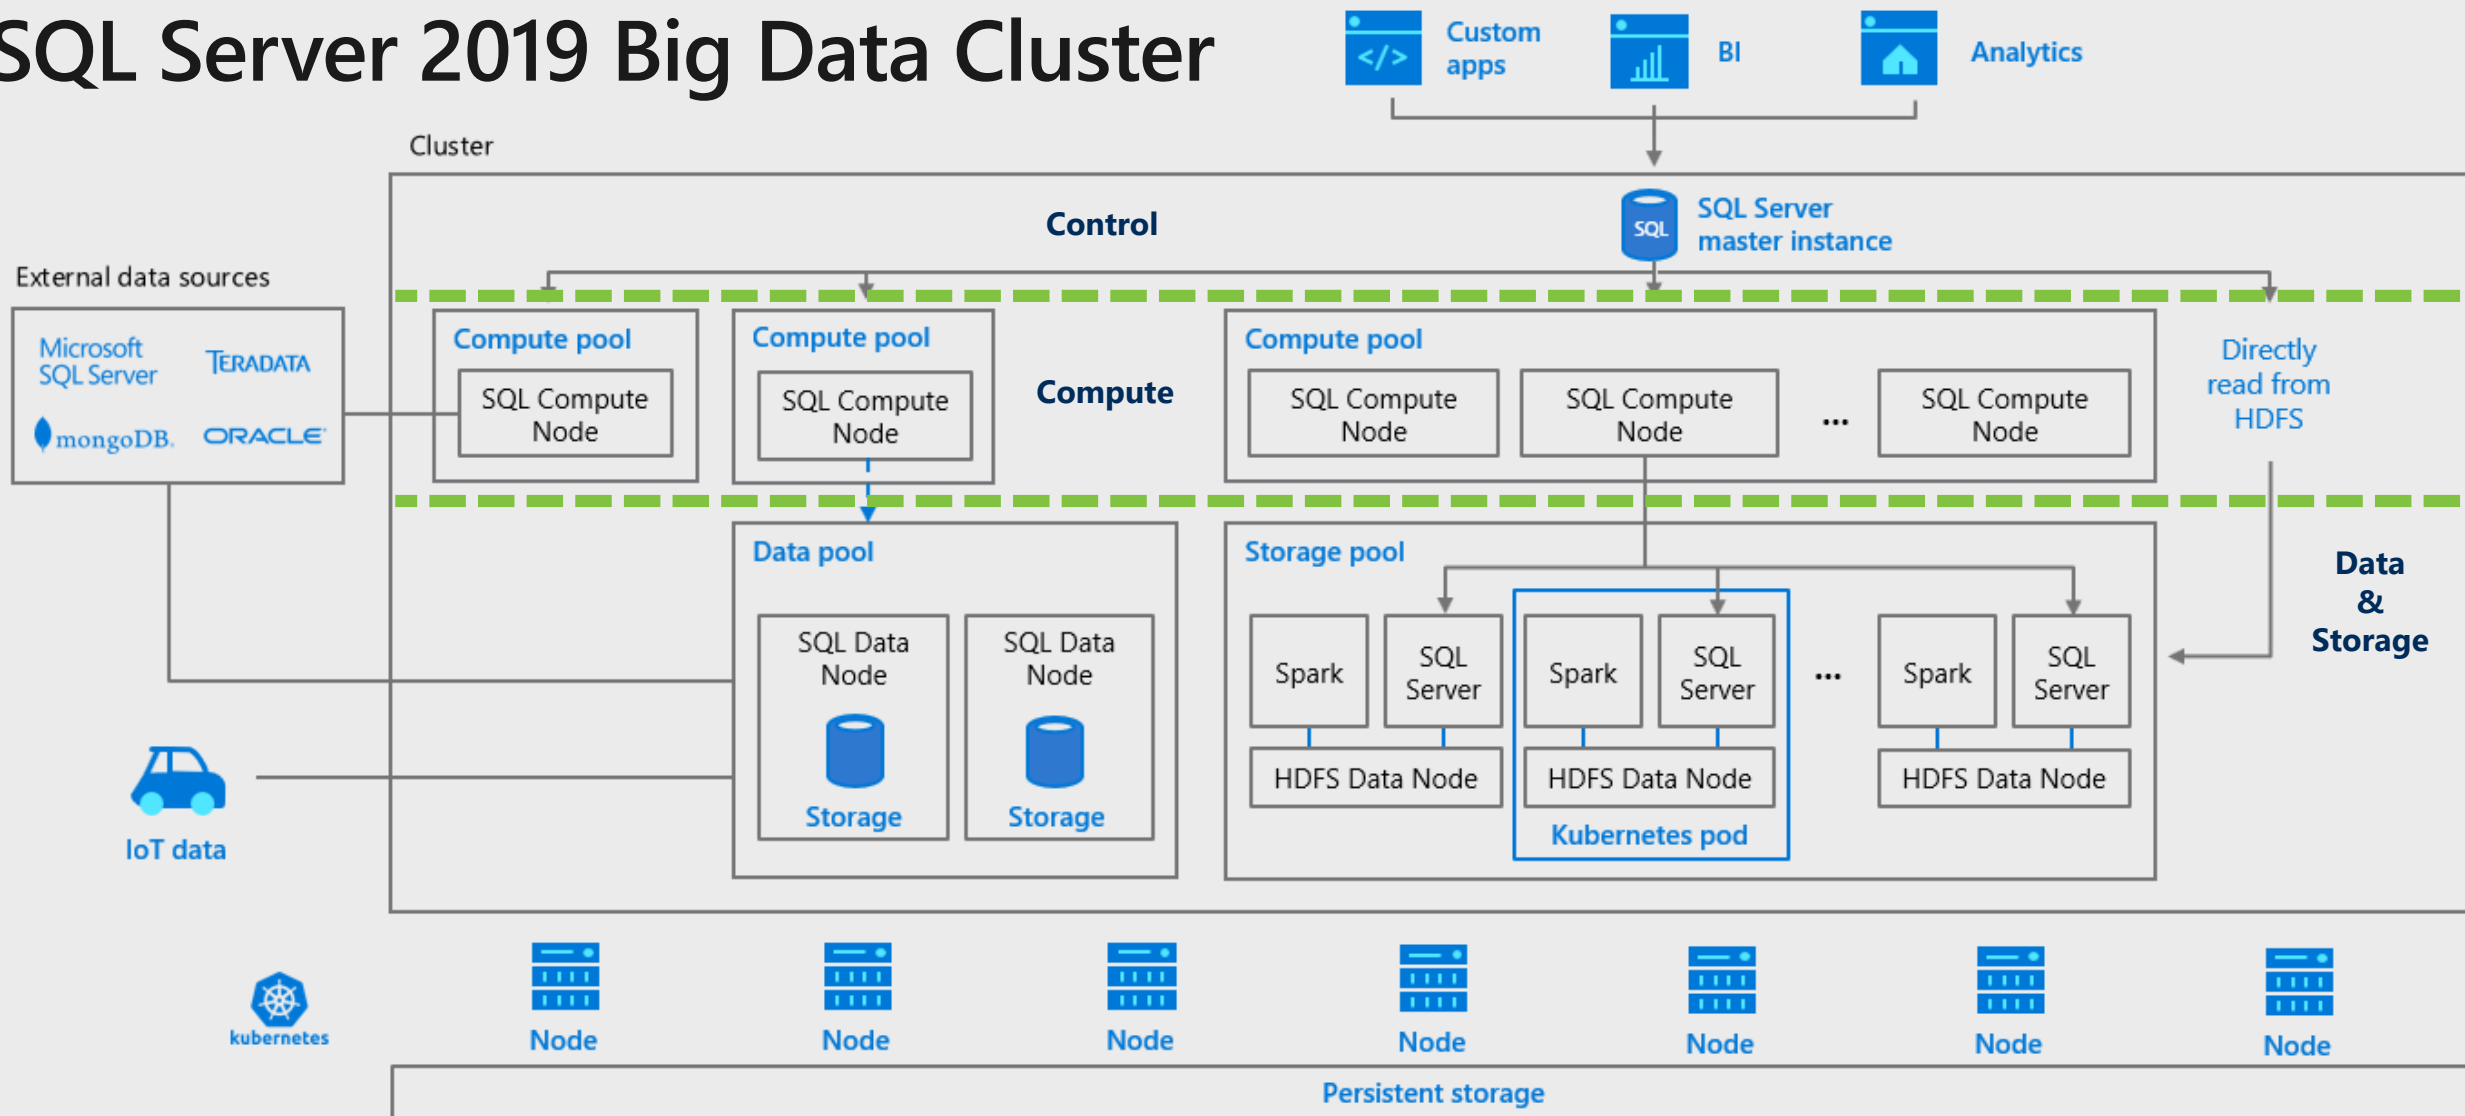
Azure Data Studio



SSMS v18



SQL Server 2019 Big Data Cluster

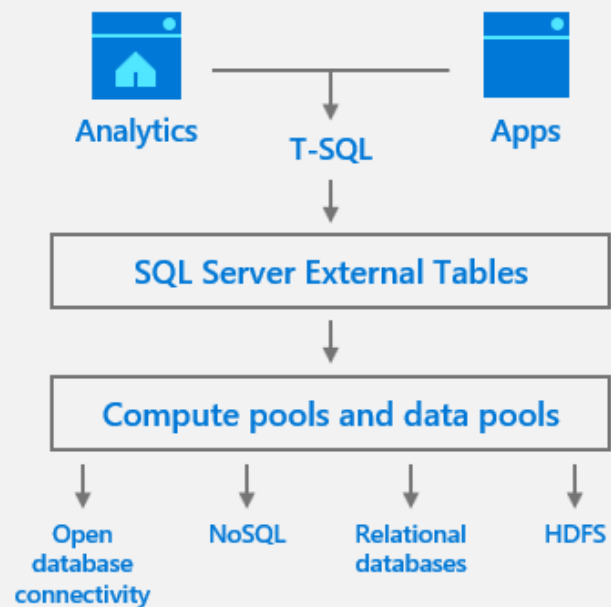


<https://cloudblogs.microsoft.com/sqlserver/2018/09/25/introducing-microsoft-sql-server-2019-big-data-clusters/>



SQL Server 2019 Big Data Cluster

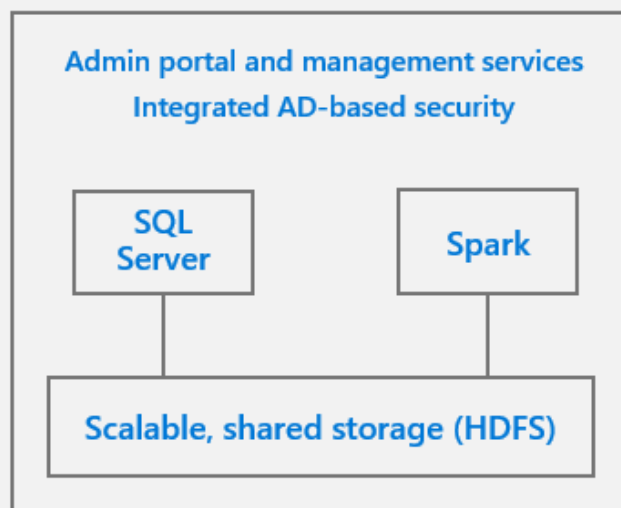
Data virtualization



Combine data from many sources without moving or replicating it

Scale out compute and caching to boost performance

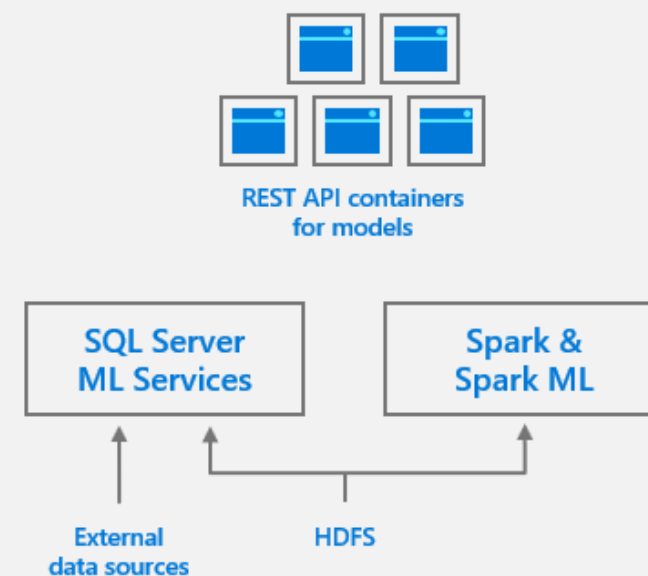
Managed SQL Server, Spark, and data lake



Store high volume data in a data lake and access it easily using either SQL or Spark

Management services, admin portal, and integrated security make it all easy to manage

Complete AI platform

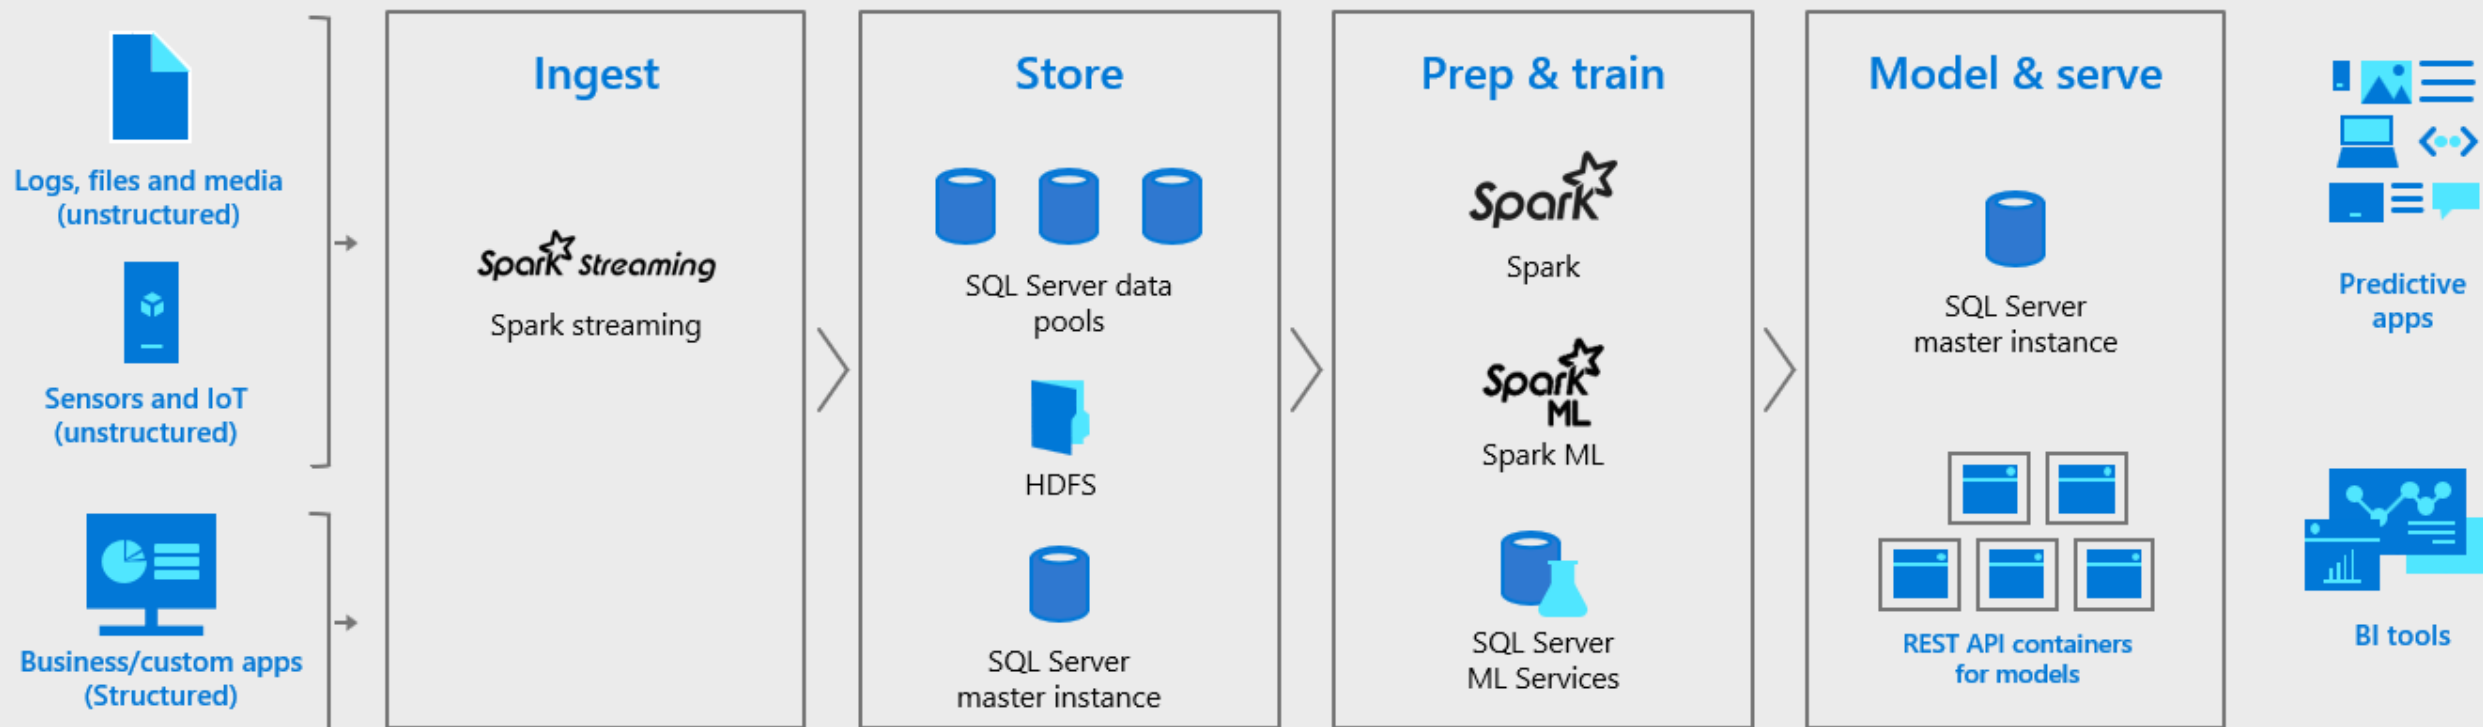


Easily feed integrated data from many sources to your model training

Ingest and prep data and then train, store, and operationalize your models all in one system

<https://cloudblogs.microsoft.com/sqlserver/2018/09/25/introducing-microsoft-sql-server-2019-big-data-clusters/>

SQL Server 2019 Big Data Cluster





<https://cloudblogs.microsoft.com/sqlserver/2018/09/25/introducing-microsoft-sql-server-2019-big-data-clusters/>

What


Why

How




What

Why




How...

do Big Data Clusters address our big data problems?
do we implement this new product architecture?




What

Why



How... do Big Data Clusters address our big data problems?
do we implement this new product architecture?





Volume



Distributed storage pools as standard.

Velocity



Containerised scale out architecture.

Veracity





Query everything in T-SQL via PolyBase.

Variety




Built in data lake capabilities via HDFS.




What

Why





How...

do Big Data Clusters address our big data problems?
do we implement this new product architecture?




SQL Server 2019 Big Data Clusters



SQL Server 2019 Big Data Clusters


kubernetes   minikube




Windows


ubuntu


Mac OS




SQL Server 2019 Big Data Clusters

kubernetes




minikube





Azure



Google Cloud Platform







SQL Server 2019 Big Data Clusters

Cloud tech, using cloud concepts, soon to be available as an on premises product.

Not only cloud.



Thanks for Listening

Paul Andrew




@MrPaulAndrew



Email: paul@mrpaulandrew.com

Blog: mrpaulandrew.com

GitHub: github.com/mrpaulandrew



Reference Material

What's new in SQL Server 2019?

<https://docs.microsoft.com/en-us/sql/sql-server/what-s-new-in-sql-server-ver15?view=sql-server-ver15>

SQL Server 2019 White Paper

<https://info.microsoft.com/ww-landing-SQLDB-Microsoft-SQL-Server-WhitePaper.html>

Introducing SQL Server 2019 Big Data Clusters – Travis Wright

<https://cloudblogs.microsoft.com/sqlserver/2018/09/25/introducing-microsoft-sql-server-2019-big-data-clusters/>

MS Ignite Session – The future of querying big data with PolyBase and SQL Server

<https://myignite.techcommunity.microsoft.com/sessions/66201>

MS Ignite Session – Deep dive on SQL Server and big data

<https://myignite.techcommunity.microsoft.com/sessions/65967>

How to deploy SQL Server big data cluster on Kubernetes

<https://docs.microsoft.com/en-us/sql/big-data-cluster/deployment-guidance?view=sql-server-ver15>

Deploy SQL Server big data cluster on Azure Kubernetes Service

<https://docs.microsoft.com/en-us/sql/big-data-cluster/quickstart-big-data-cluster-deploy>

