

Azure Stream Analytics

Beyond IoT Real-time Data Ingestion



Paul Andrew



@MrPaulAndrew | mrpaulandrew.com

Session



intelligent
cloud conference

Event
partners



Expo
partners



Expo light
partners



Azure Stream Analytics

Beyond IoT Real-time Data Ingestion



Paul Andrew



@MrPaulAndrew | mrpaulandrew.com

Session



intelligent
cloud conference



intelligent
cloud conference

Azure Stream Analytics



Real-time data problems

What is ASA and why use it

Production Considerations

Lambda Architecture

Azure Stream Analytics



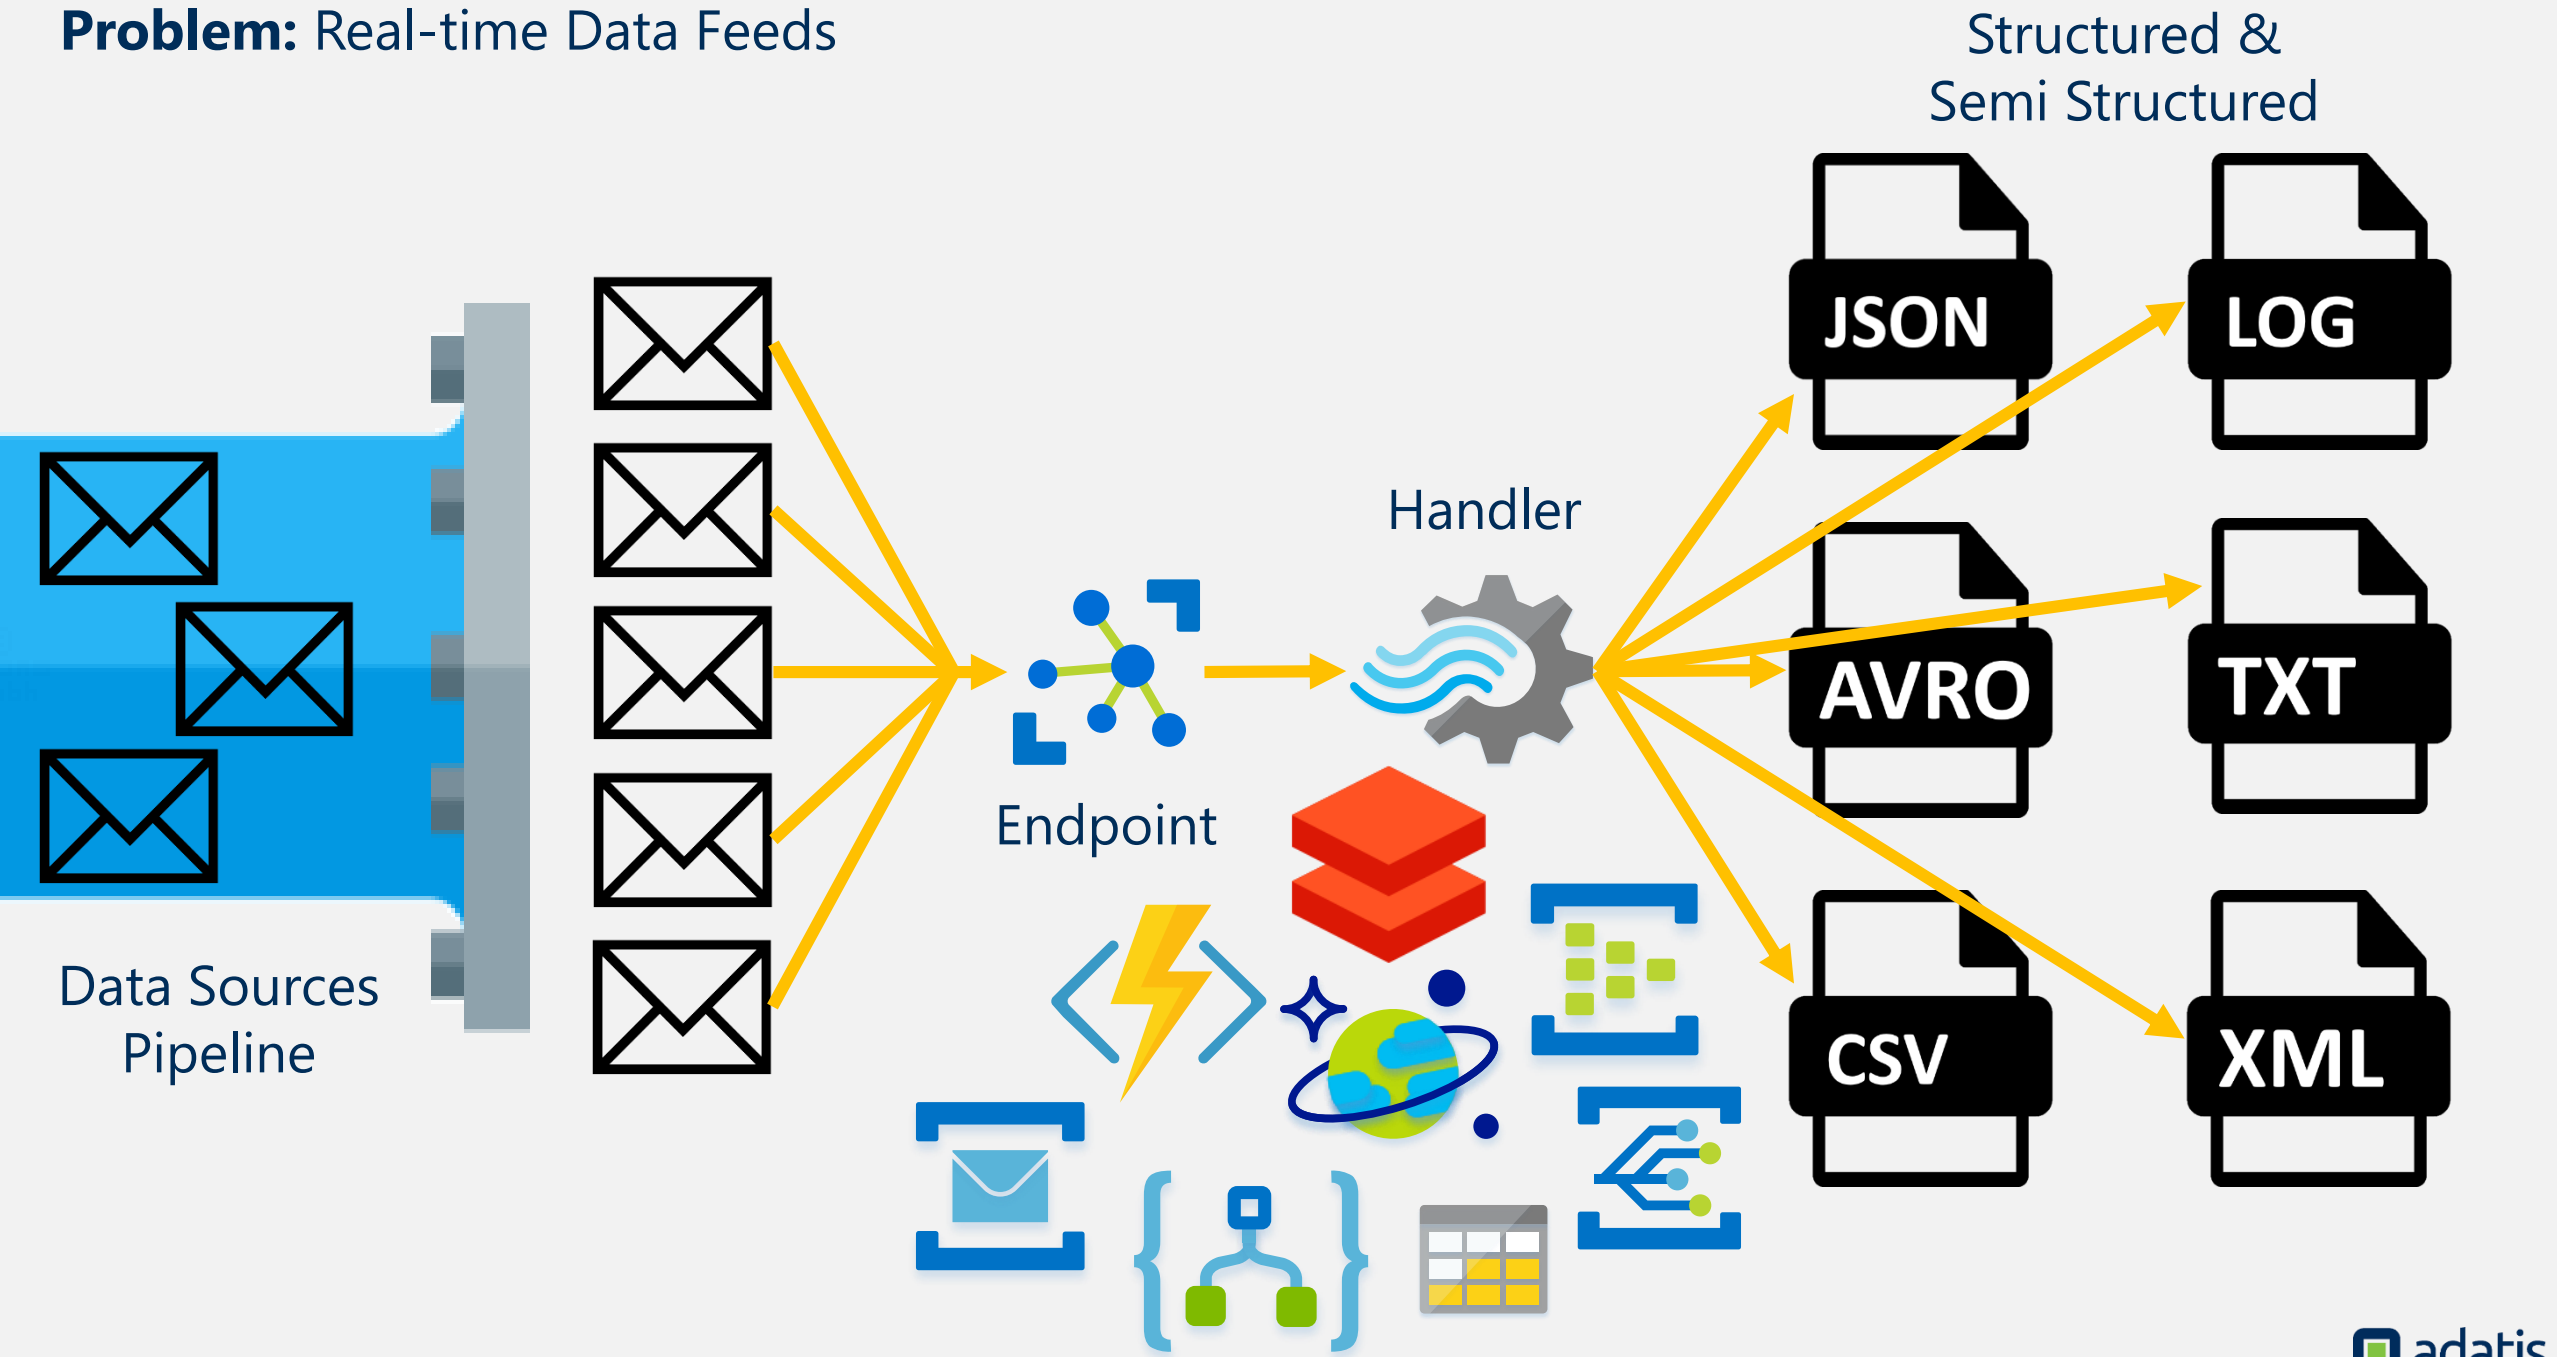
Real-time data problems

What is *ASA* and why use it

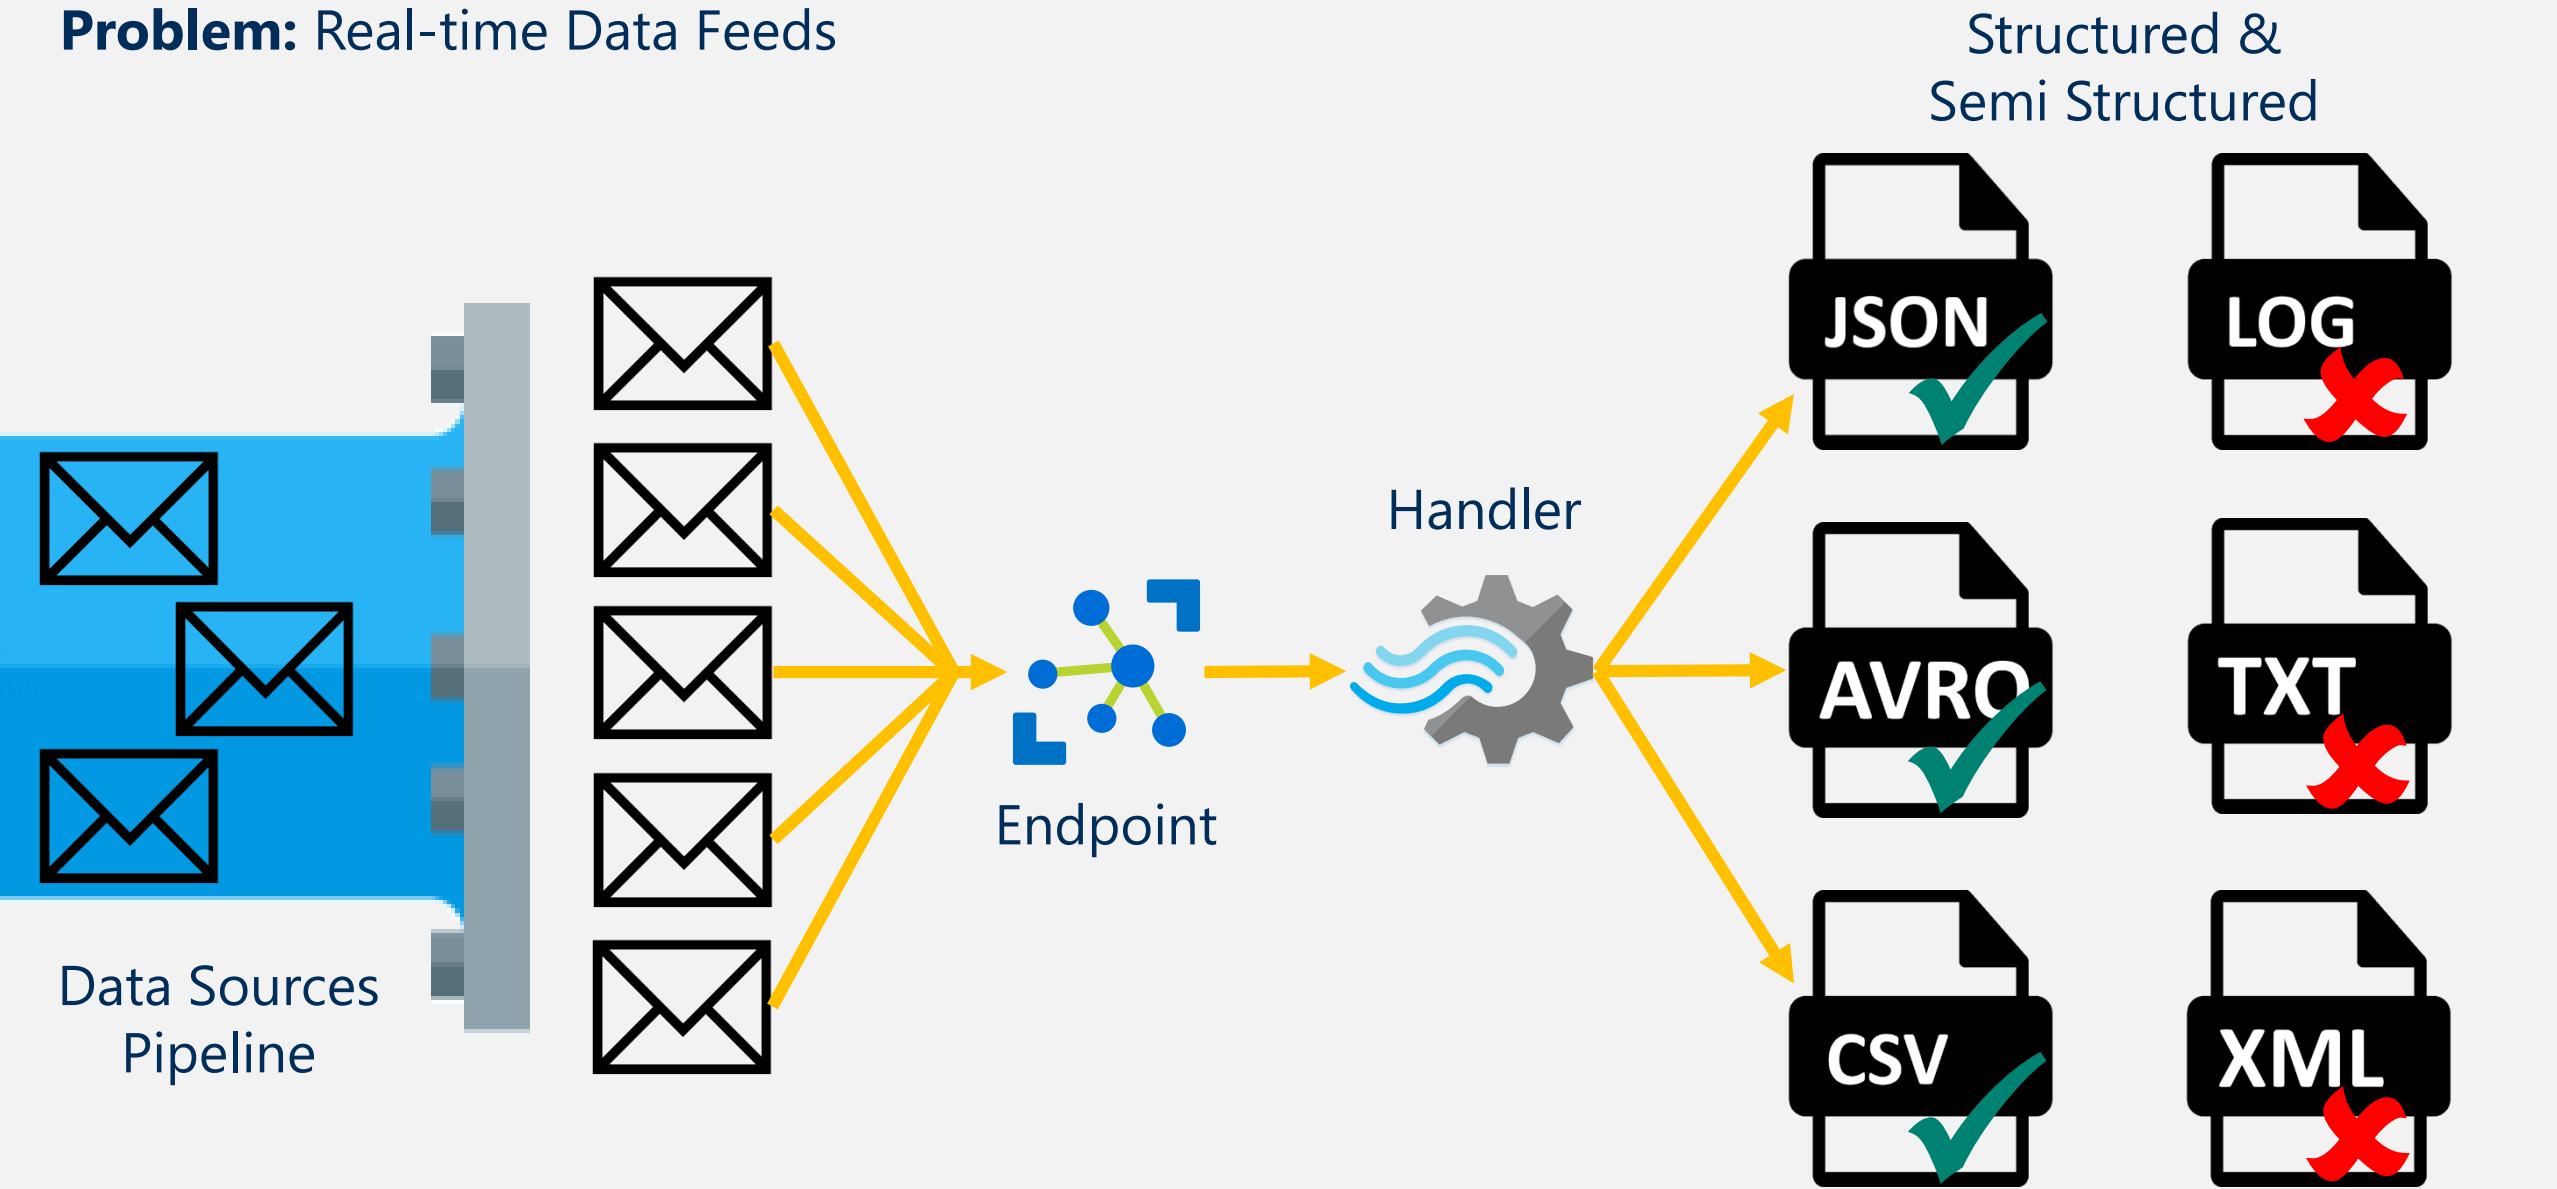
Production Considerations

Lambda Architecture

Problem: Real-time Data Feeds



Problem: Real-time Data Feeds



Azure Stream Analytics



Real-time data problems

What is ASA and why use it

Production Considerations

Lambda Architecture

Microsoft Azure

Contact Sales: 0800-026-0797

Search

Overview ▾ Solutions **Products ▾** Documentation Pricing Training Marketplace ▾ Partners ▾ Support ▾ Blog More ▾

Azure Stream Analytics

An on-demand real-time analytics service to power intelligent action

Easily develop and run massively parallel real-time analytics on multiple IoT or non-IoT streams of data using simple SQL-like language. Use customised code for advanced scenarios. With no infrastructure to manage, you can process data on demand, scale instantly and only pay per job.

Get started for free >

Azure Stream Analytics



<https://azure.microsoft.com/en-gb/services/stream-analytics/>

Real-time data problems

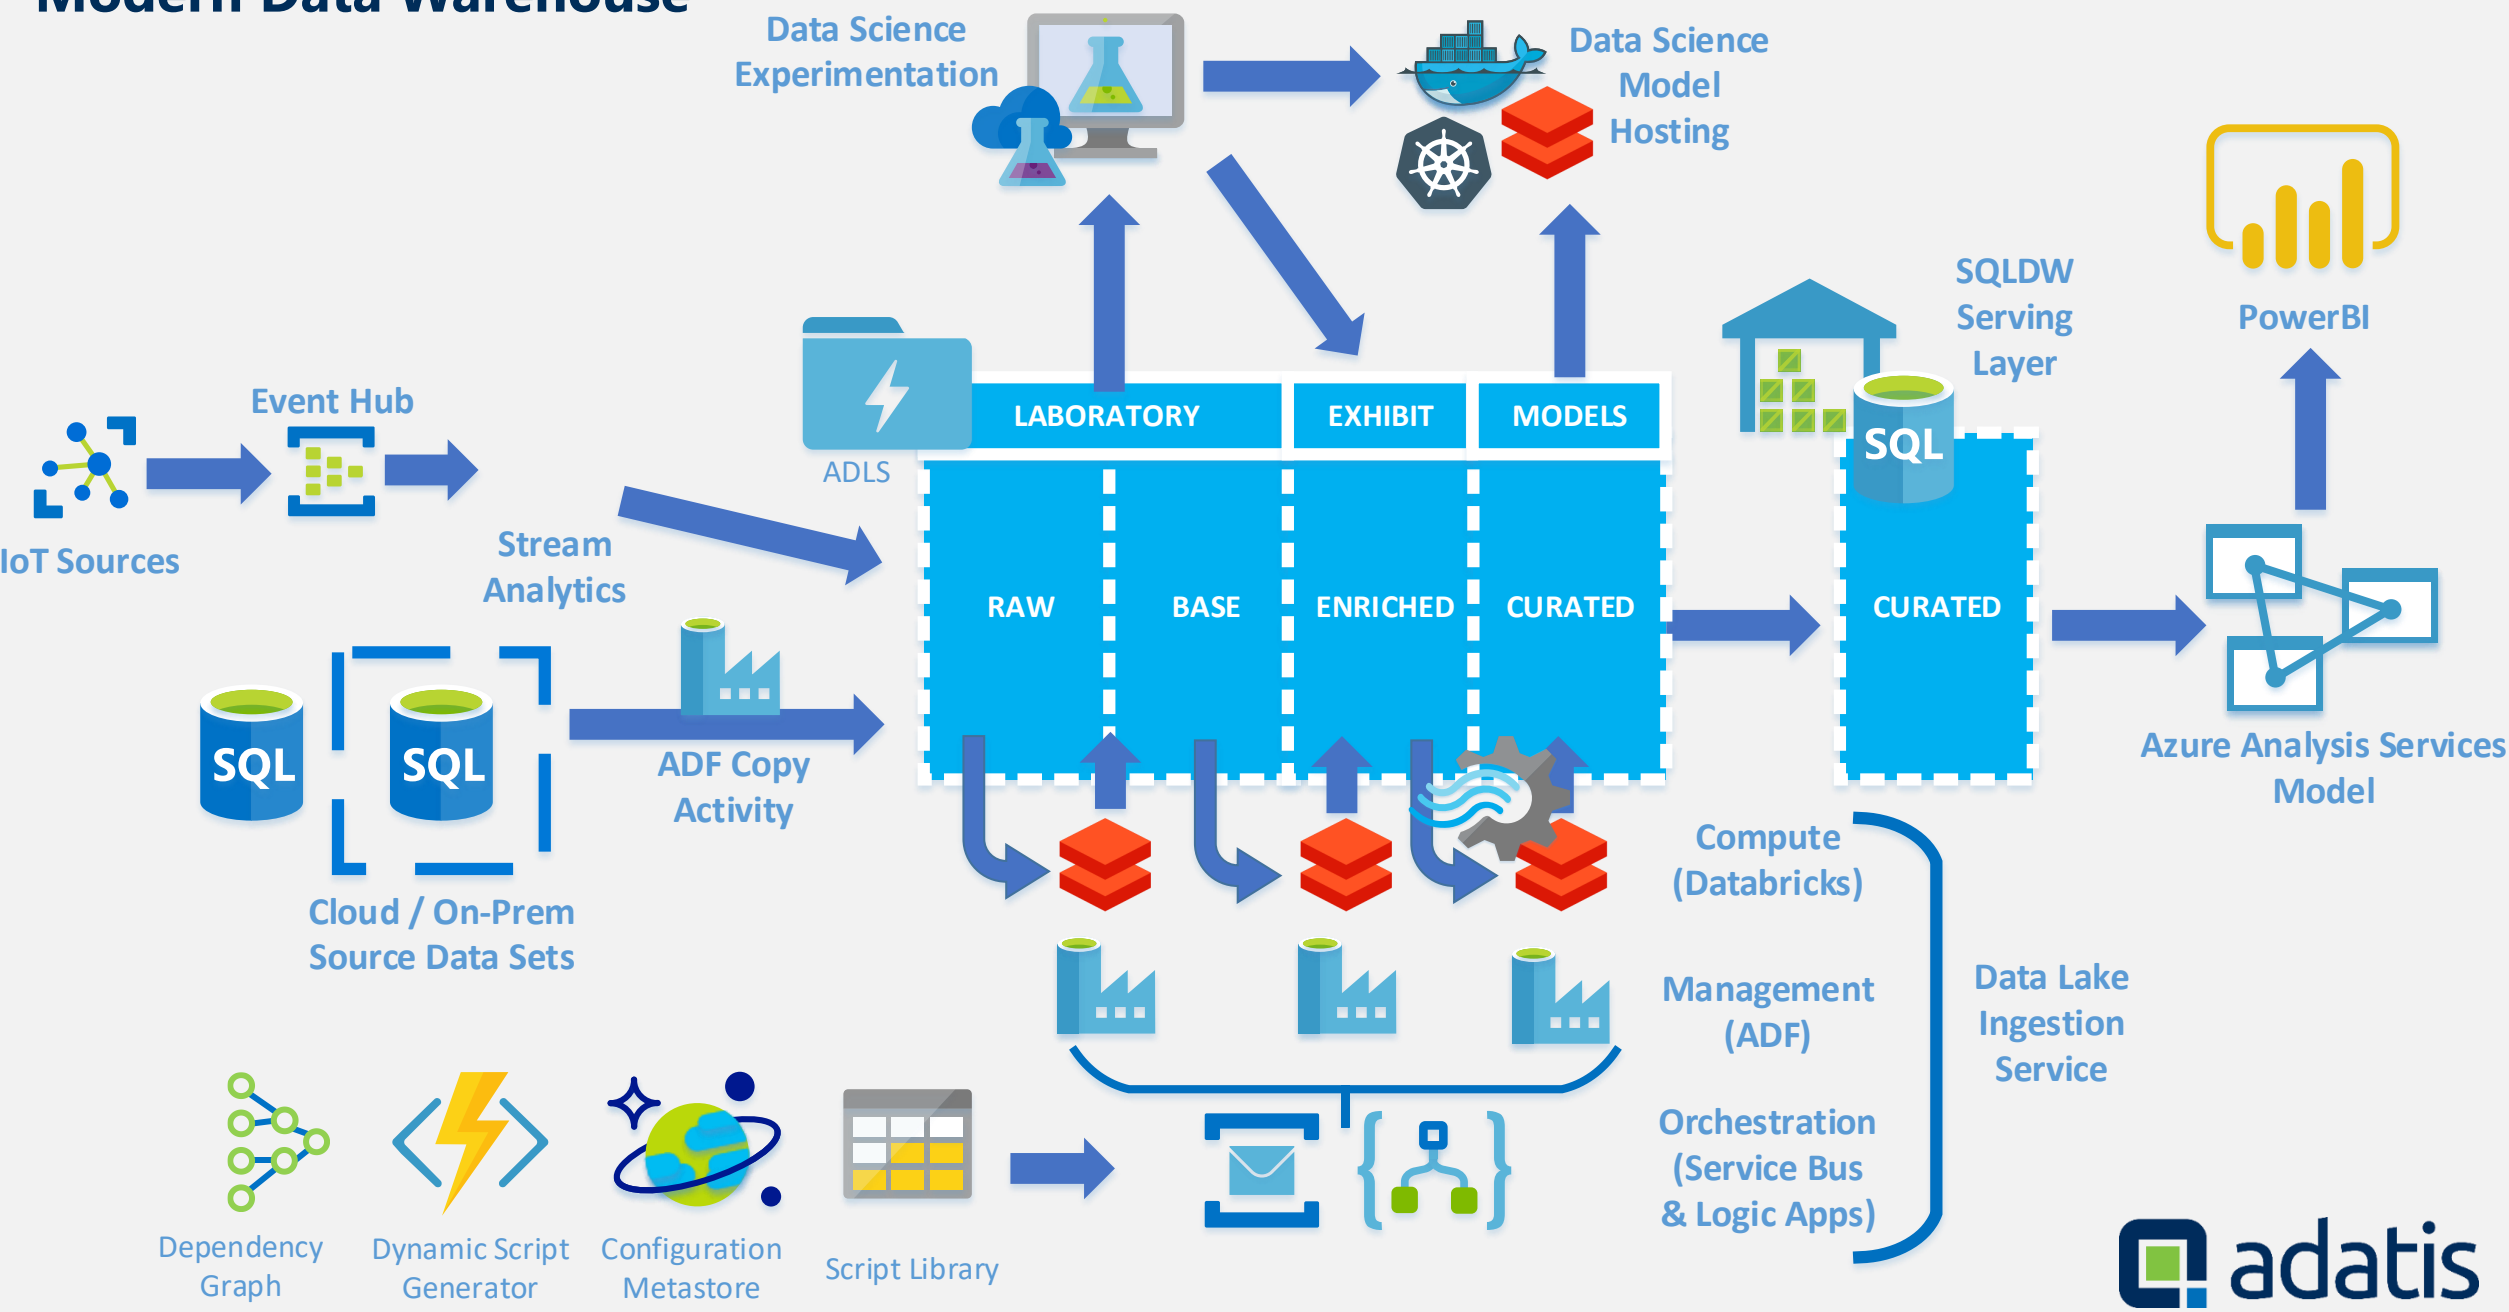
What is ASA and why use it

Production Considerations


Lambda Architecture




Modern Data Warehouse



Modern Data Warehouse




Azure IoT Hub vs Azure Event Hub



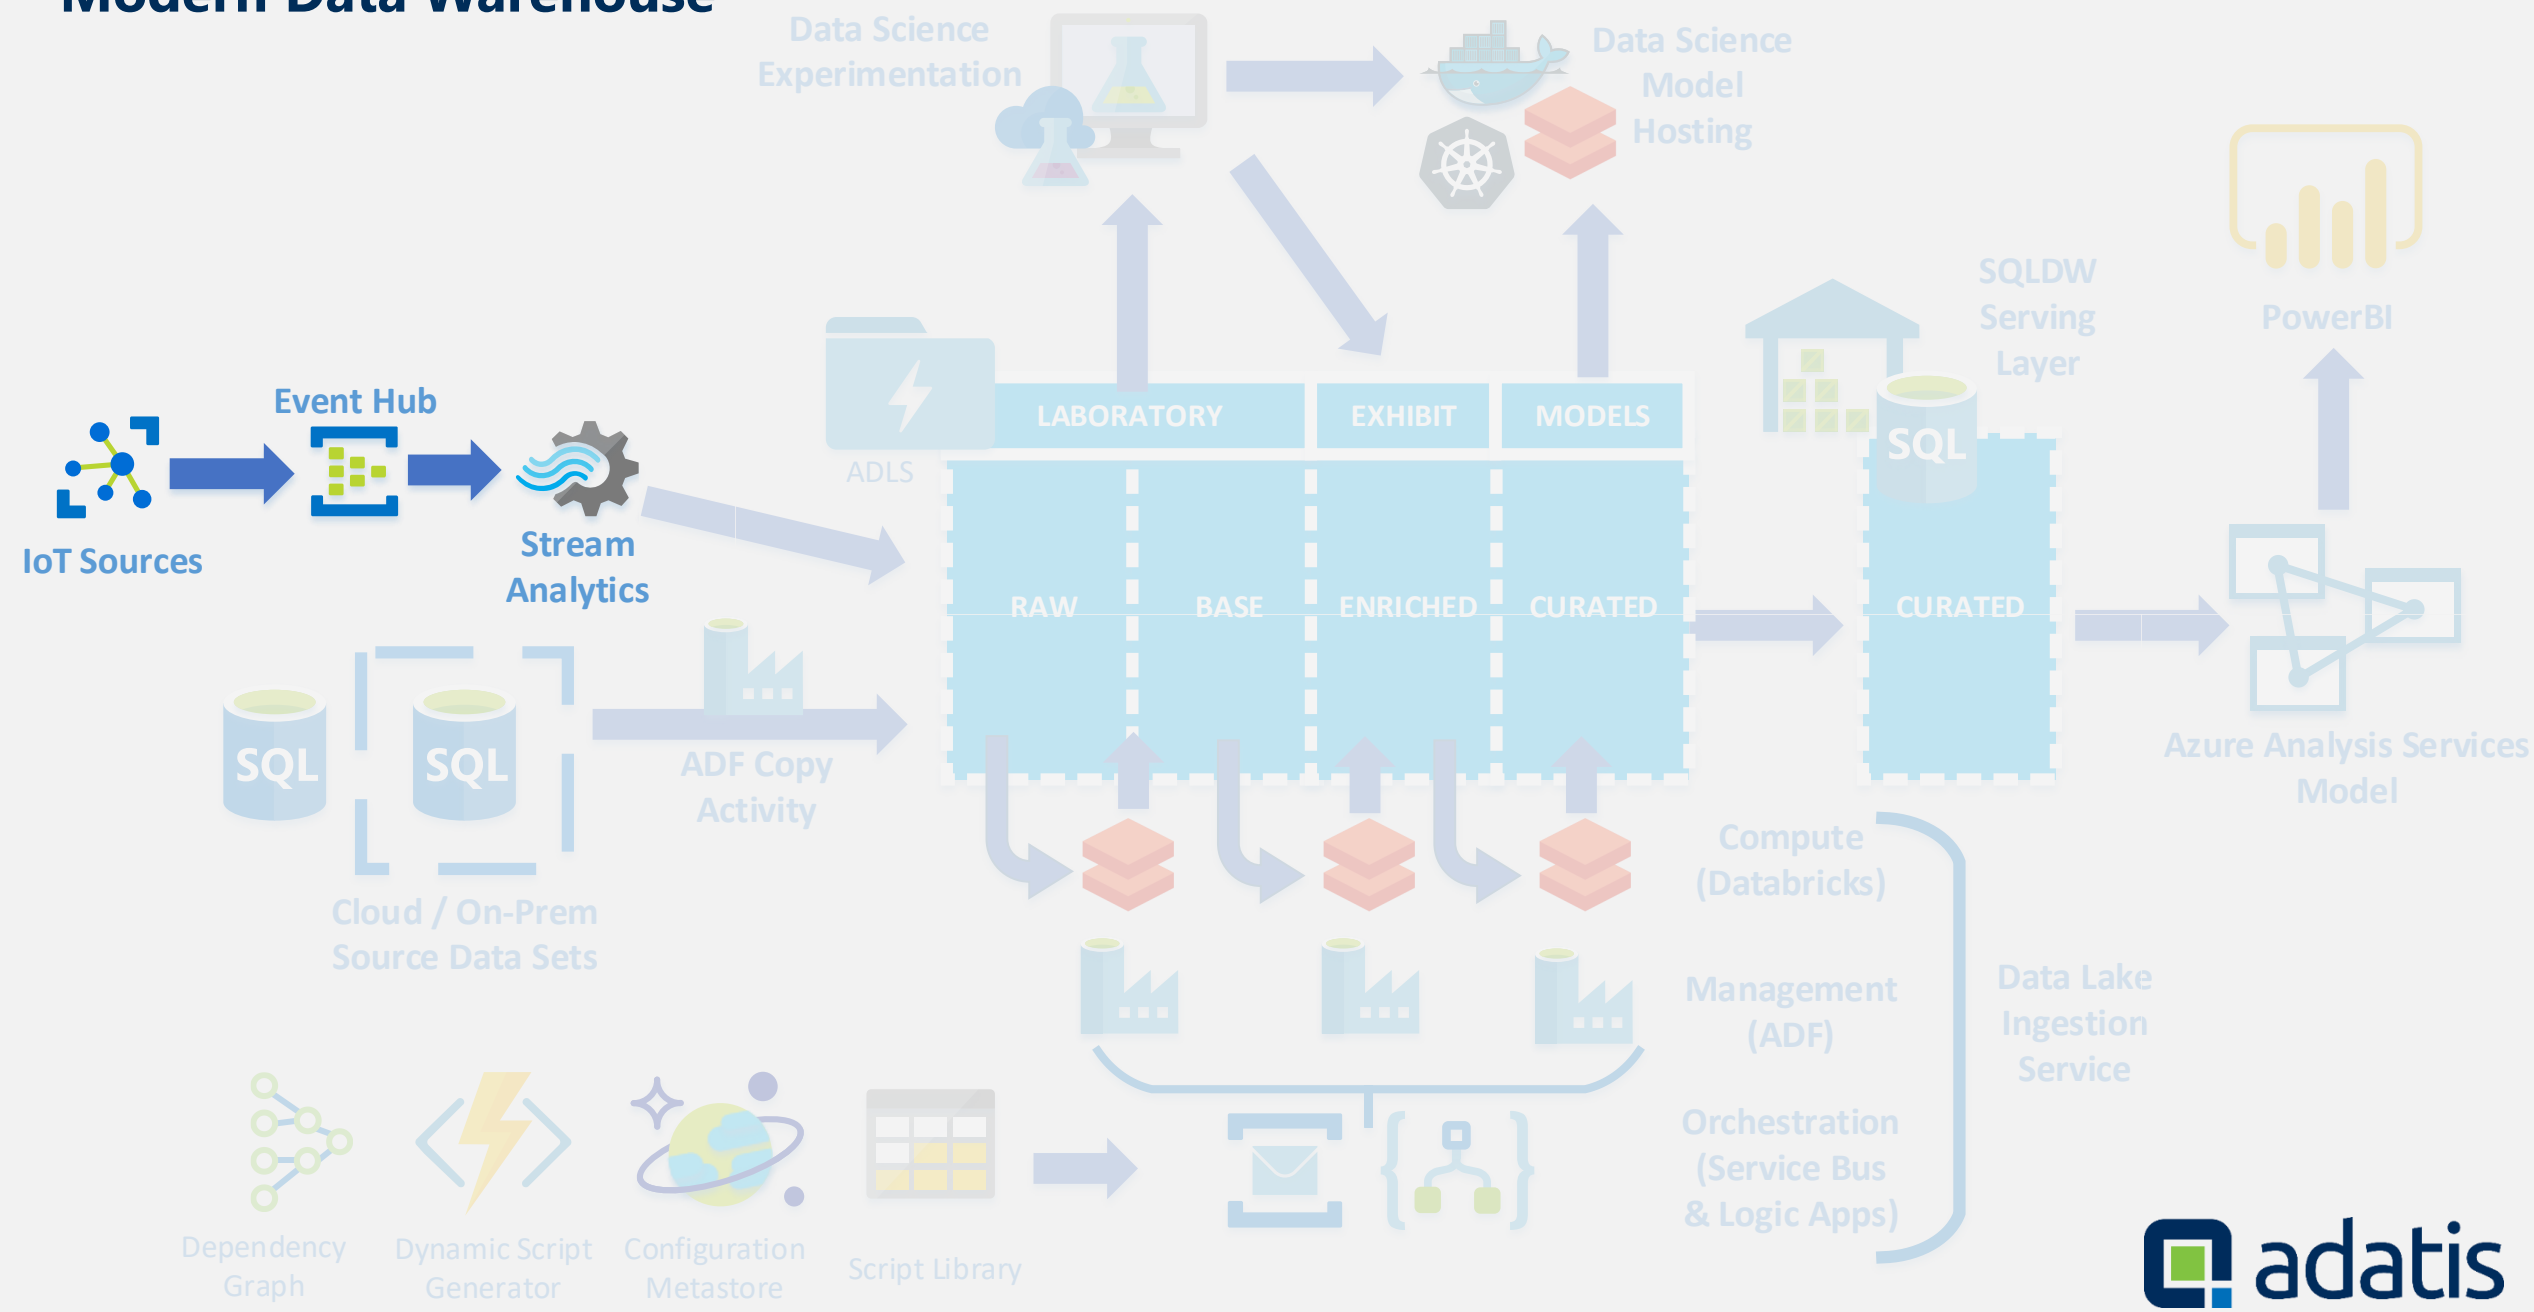
Feature	Azure IoT Hub	Azure Event Hub
Message Direction	2 Way	1 Way
Protocol Support	MQTT, AMQP, HTTP	AMQP, HTTP
Scaling	Configured	Automatic
Message Routing	Yes	No
Security	Device Level	Hub Level
Device State Support	Yes	No
Message Capturing	No	Yes
Multiple Namespaces	No	Yes
Tiers	F1/S1/S2/S3	Basic/Standard
Service Endpoint	Yes	Yes (preview)

Azure IoT Hub vs Azure Event Hub



Feature	Azure IoT Hub	Azure Event Hub
Message Direction	2 Way	1 Way
Protocol Support	MQTT, AMQP, HTTP	AMQP, HTTP
Scaling	Configured	Automatic
Message Routing	Yes	No
Security	Device Level	Hub Level
Device State Support	Yes	No
Message Capturing	No	Yes
Multiple Namespaces	No	Yes
Tiers	F1/S1/S2/S3	Basic/Standard
Service Endpoint	Yes	Yes (preview)

Modern Data Warehouse




Azure Stream Analytics



Inputs **Outputs**



Azure Stream Analytics



Azure Stream Analytics

Query

```
SELECT  
  
    SUM(CAST(eh.UnitPrice AS float)) AS UnitPrice,  
    SUM(CAST(eh.LineTotal AS float)) AS LineTotal,  
    SUM(CAST(eh.OrderQty AS float)) AS OrderQty,  
    COUNT(*) AS RecordCount  
INTO  
    [powerbi]  
FROM  
    [eventhub] AS eh  
  
GROUP BY  
    eh.EventEnqueuedUtcTime,  
    SlidingWindow(second, 30)
```



Inputs



Outputs



Azure Stream Analytics

Query

```
SELECT
    eh.EventEnqueuedUtcTime,
    prd.Name AS ProductName,
    SUM(CAST(eh.UnitPrice AS float)) AS UnitPrice,
    SUM(CAST(eh.LineTotal AS float)) AS LineTotal,
    SUM(CAST(eh.OrderQty AS float)) AS OrderQty,
    COUNT(*) AS RecordCount
INTO
    [powerbi]
FROM
    [eventhub] AS eh
    INNER JOIN [Products] AS prd
        ON eh.[ProductId] = prd.[ProductId]
GROUP BY
    eh.EventEnqueuedUtcTime,
    prd.Name,
    SlidingWindow(second, 30)
```



*100MB
Limit

Inputs



Outputs




Azure Stream Analytics

Query

```
SELECT
    eh.EventEnqueuedUtcTime,
    udf.CleanString(prd.Name) AS ProductName,
    SUM(CAST(eh.UnitPrice AS float)) AS UnitPrice,
    SUM(CAST(eh.LineTotal AS float)) AS LineTotal,
    SUM(CAST(eh.OrderQty AS float)) AS OrderQty,
    COUNT(*) AS RecordCount
INTO
    [powerbi]
FROM
    [eventhub] AS eh
    INNER JOIN [Products] AS prd
        ON eh.[ProductId] = prd.[ProductId]
GROUP BY
    eh.EventEnqueuedUtcTime,
    prd.Name,
    SlidingWindow(second, 30)
```



*100MB
Limit




Inputs... Source Data
Reference Data
Custom Functions

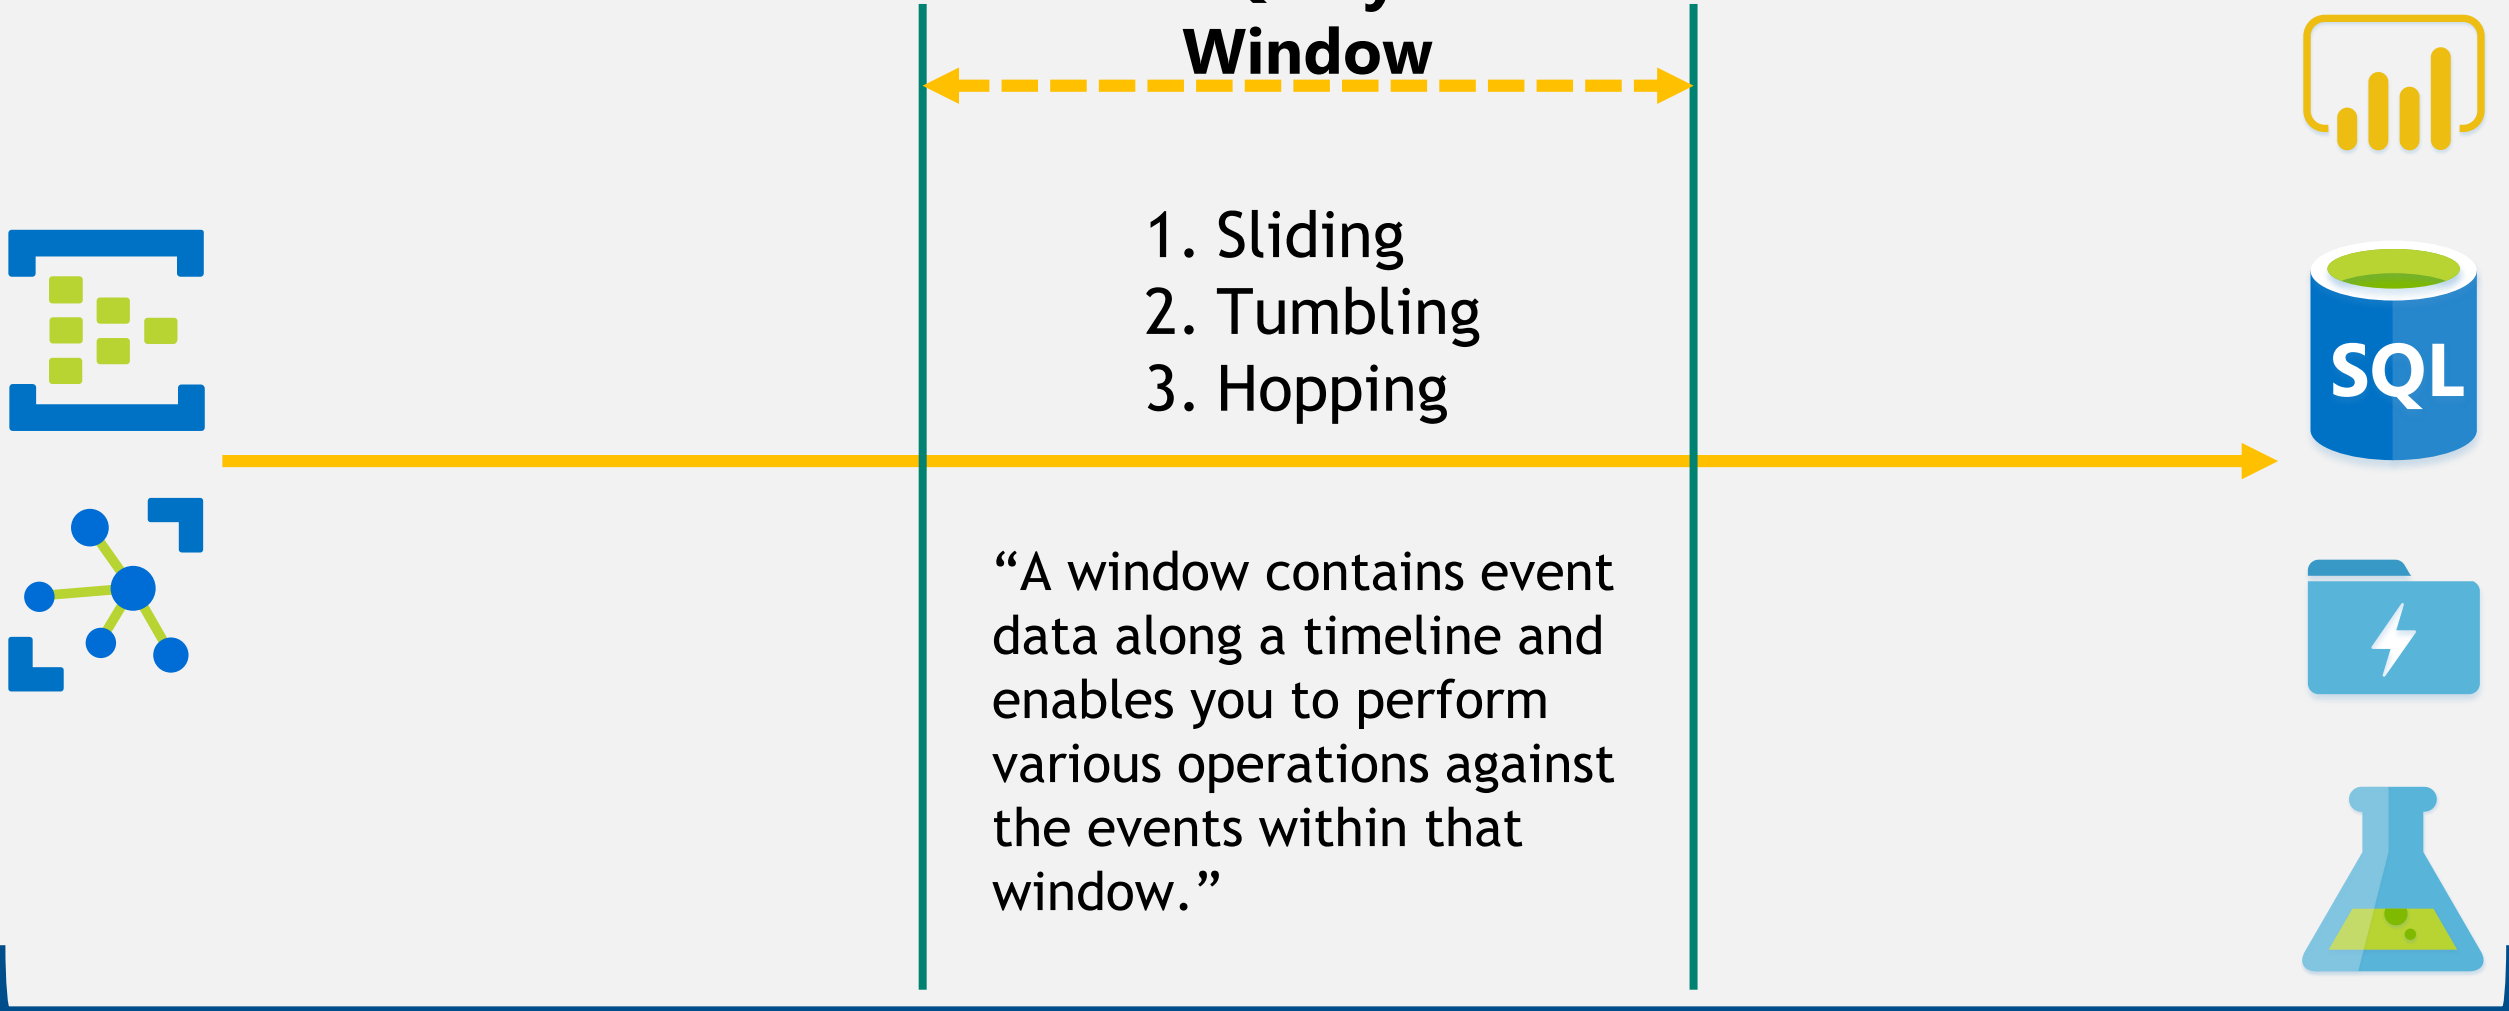


Outputs

Azure Stream Analytics

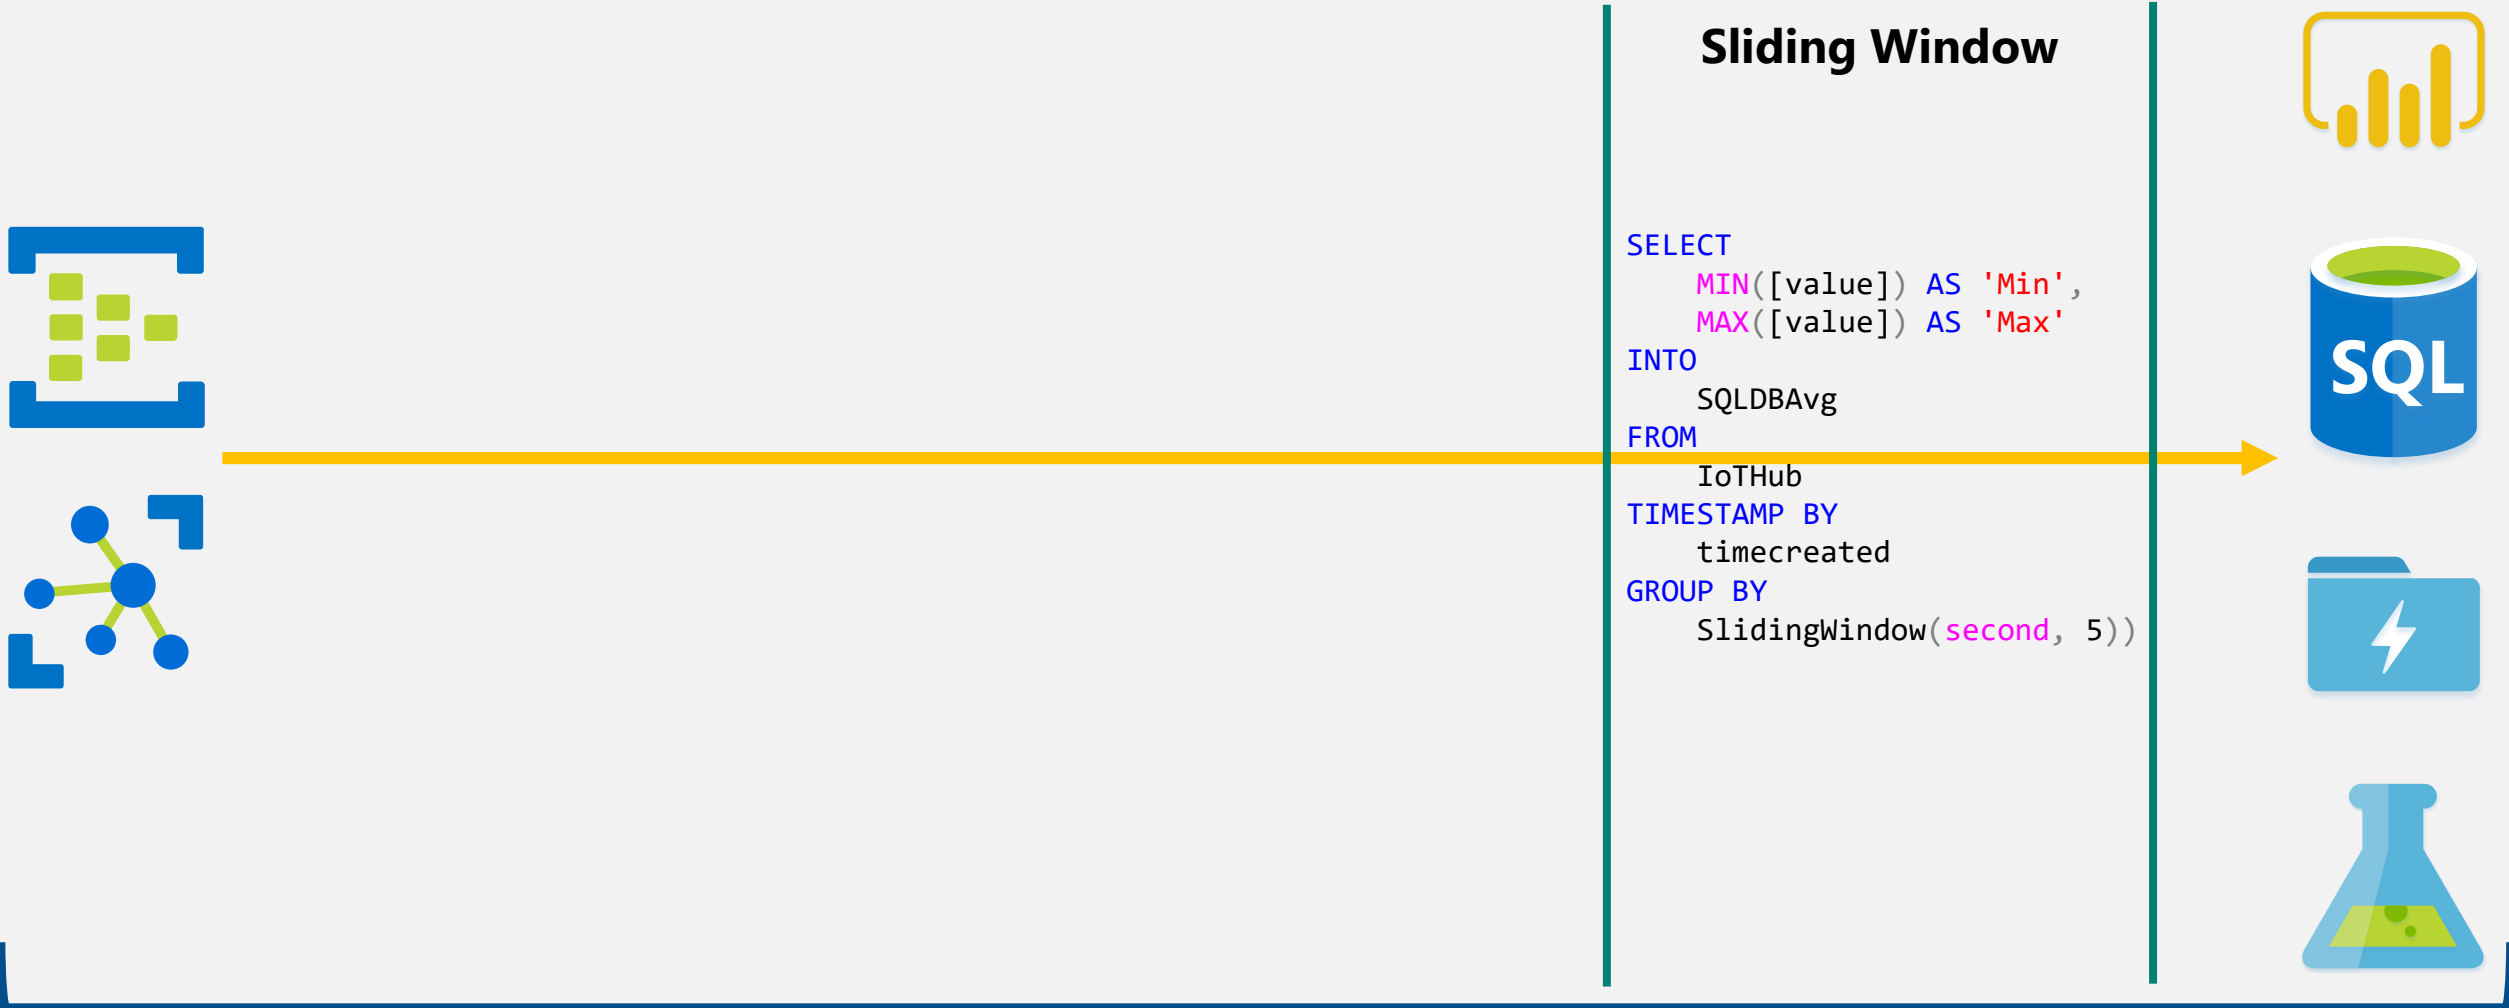


Azure Stream Analytics



Inputs **Outputs**

Azure Stream Analytics




Sliding Window

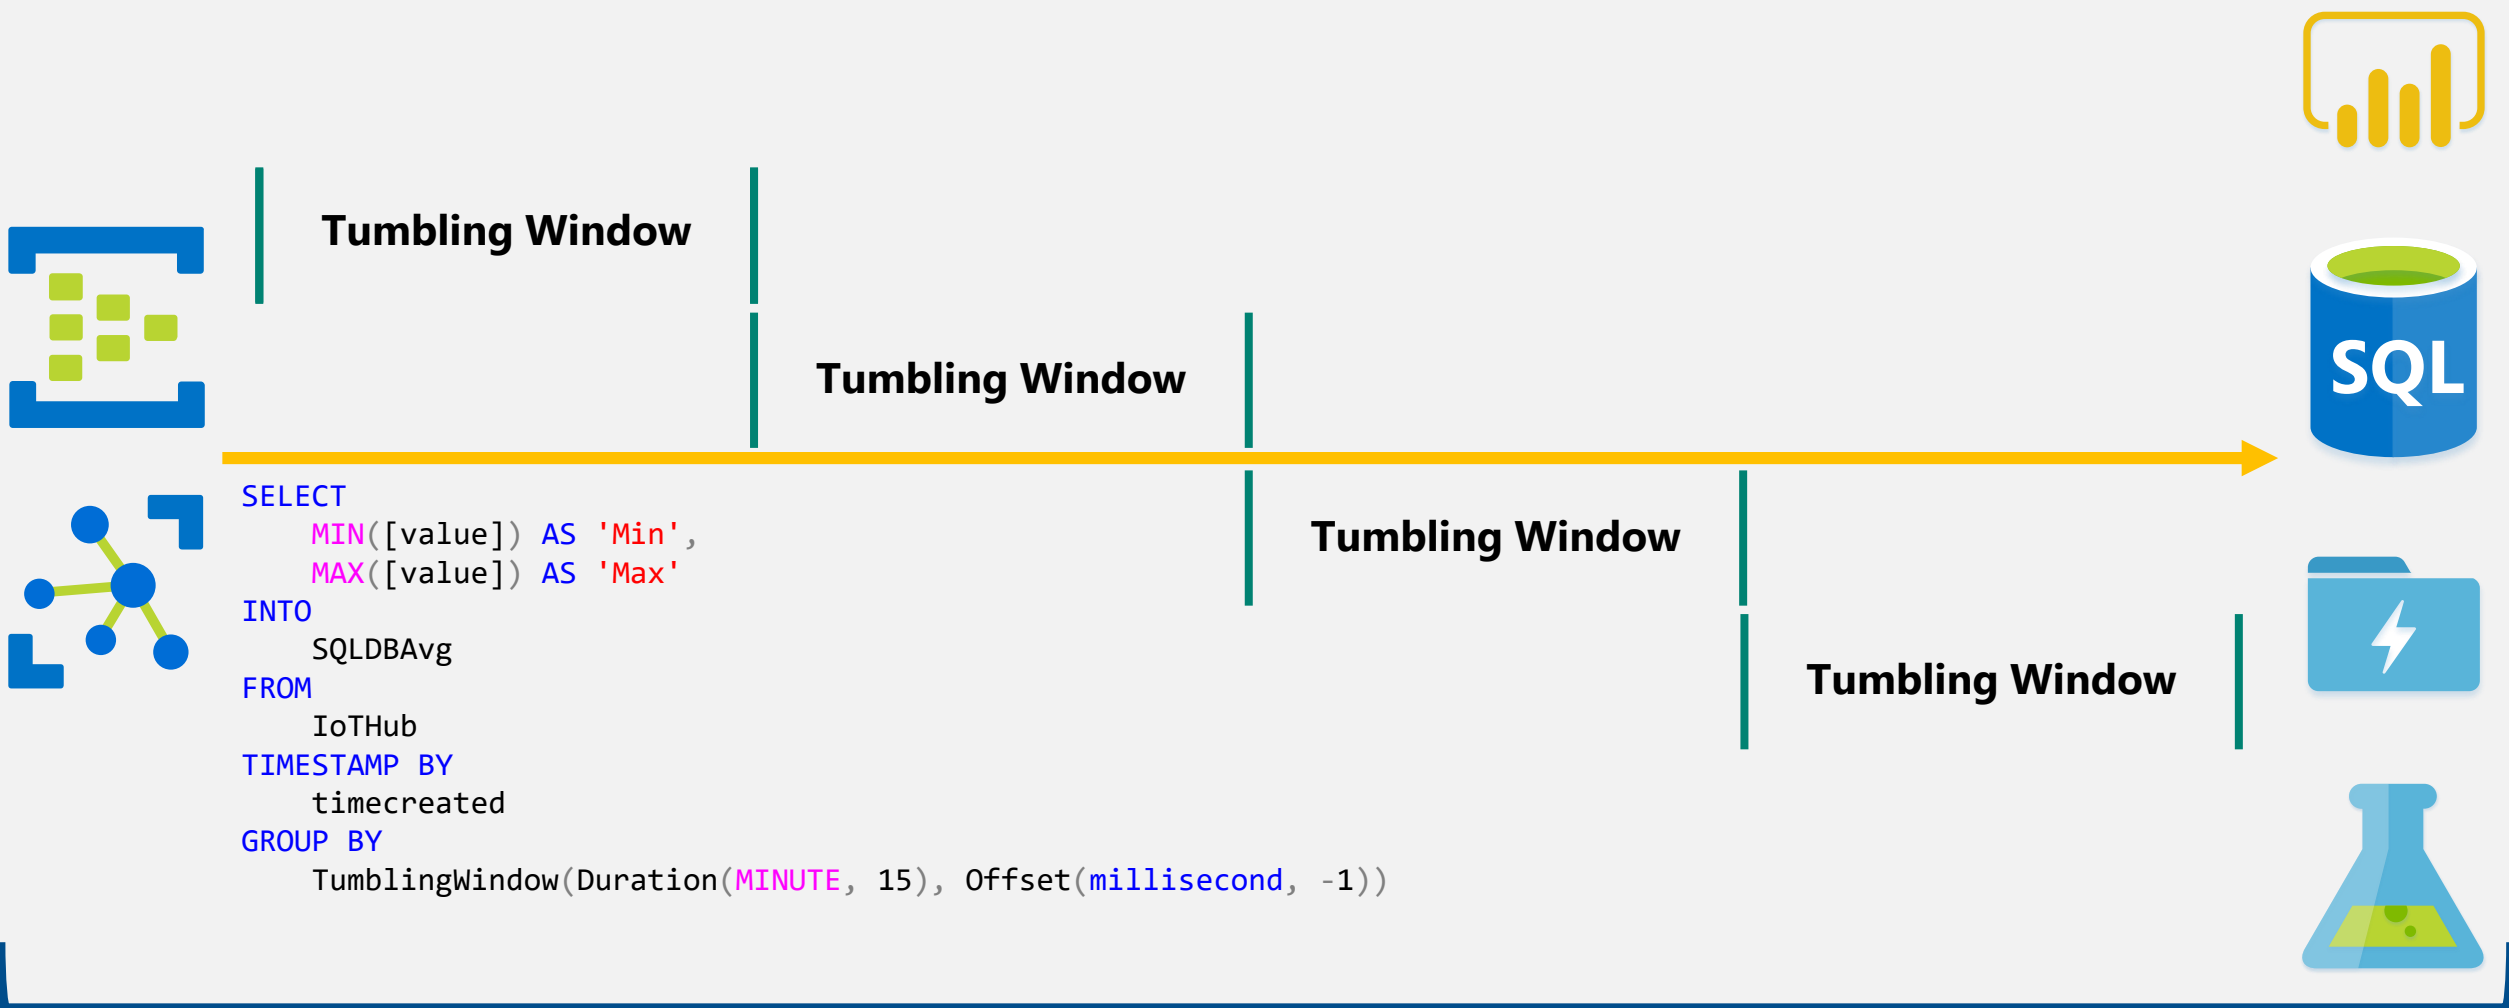
```
SELECT
    MIN([value]) AS 'Min',
    MAX([value]) AS 'Max'
INTO
    SQLDBAvg
FROM
    IoTHub
TIMESTAMP BY
    timecreated
GROUP BY
    SlidingWindow(second, 5))
```

InputsOutputs

Azure Stream Analytics



Azure Stream Analytics



```
SELECT
    MIN([value]) AS 'Min',
    MAX([value]) AS 'Max'
INTO
    SQLDBAvg
FROM
    IoTHub
TIMESTAMP BY
    timecreated
GROUP BY
    TumblingWindow(Duration(MINUTE, 15), Offset(milliseconds, -1))
```

Inputs

Outputs


Azure Stream Analytics



Inputs **Outputs**



Azure Stream Analytics



```
SELECT
    MIN([value]) AS 'Min',
    MAX([value]) AS 'Max'
INTO
    SQLDBAvg
FROM
    IoTHub
TIMESTAMP BY
    timecreated
GROUP BY
    HoppingWindow(Duration(MINUTE, 15), Hop(MINUTE, 5), Offset(milliseconds, -1))
```

Hopping Window

Hopping Window

Hopping Window

Hopping Window

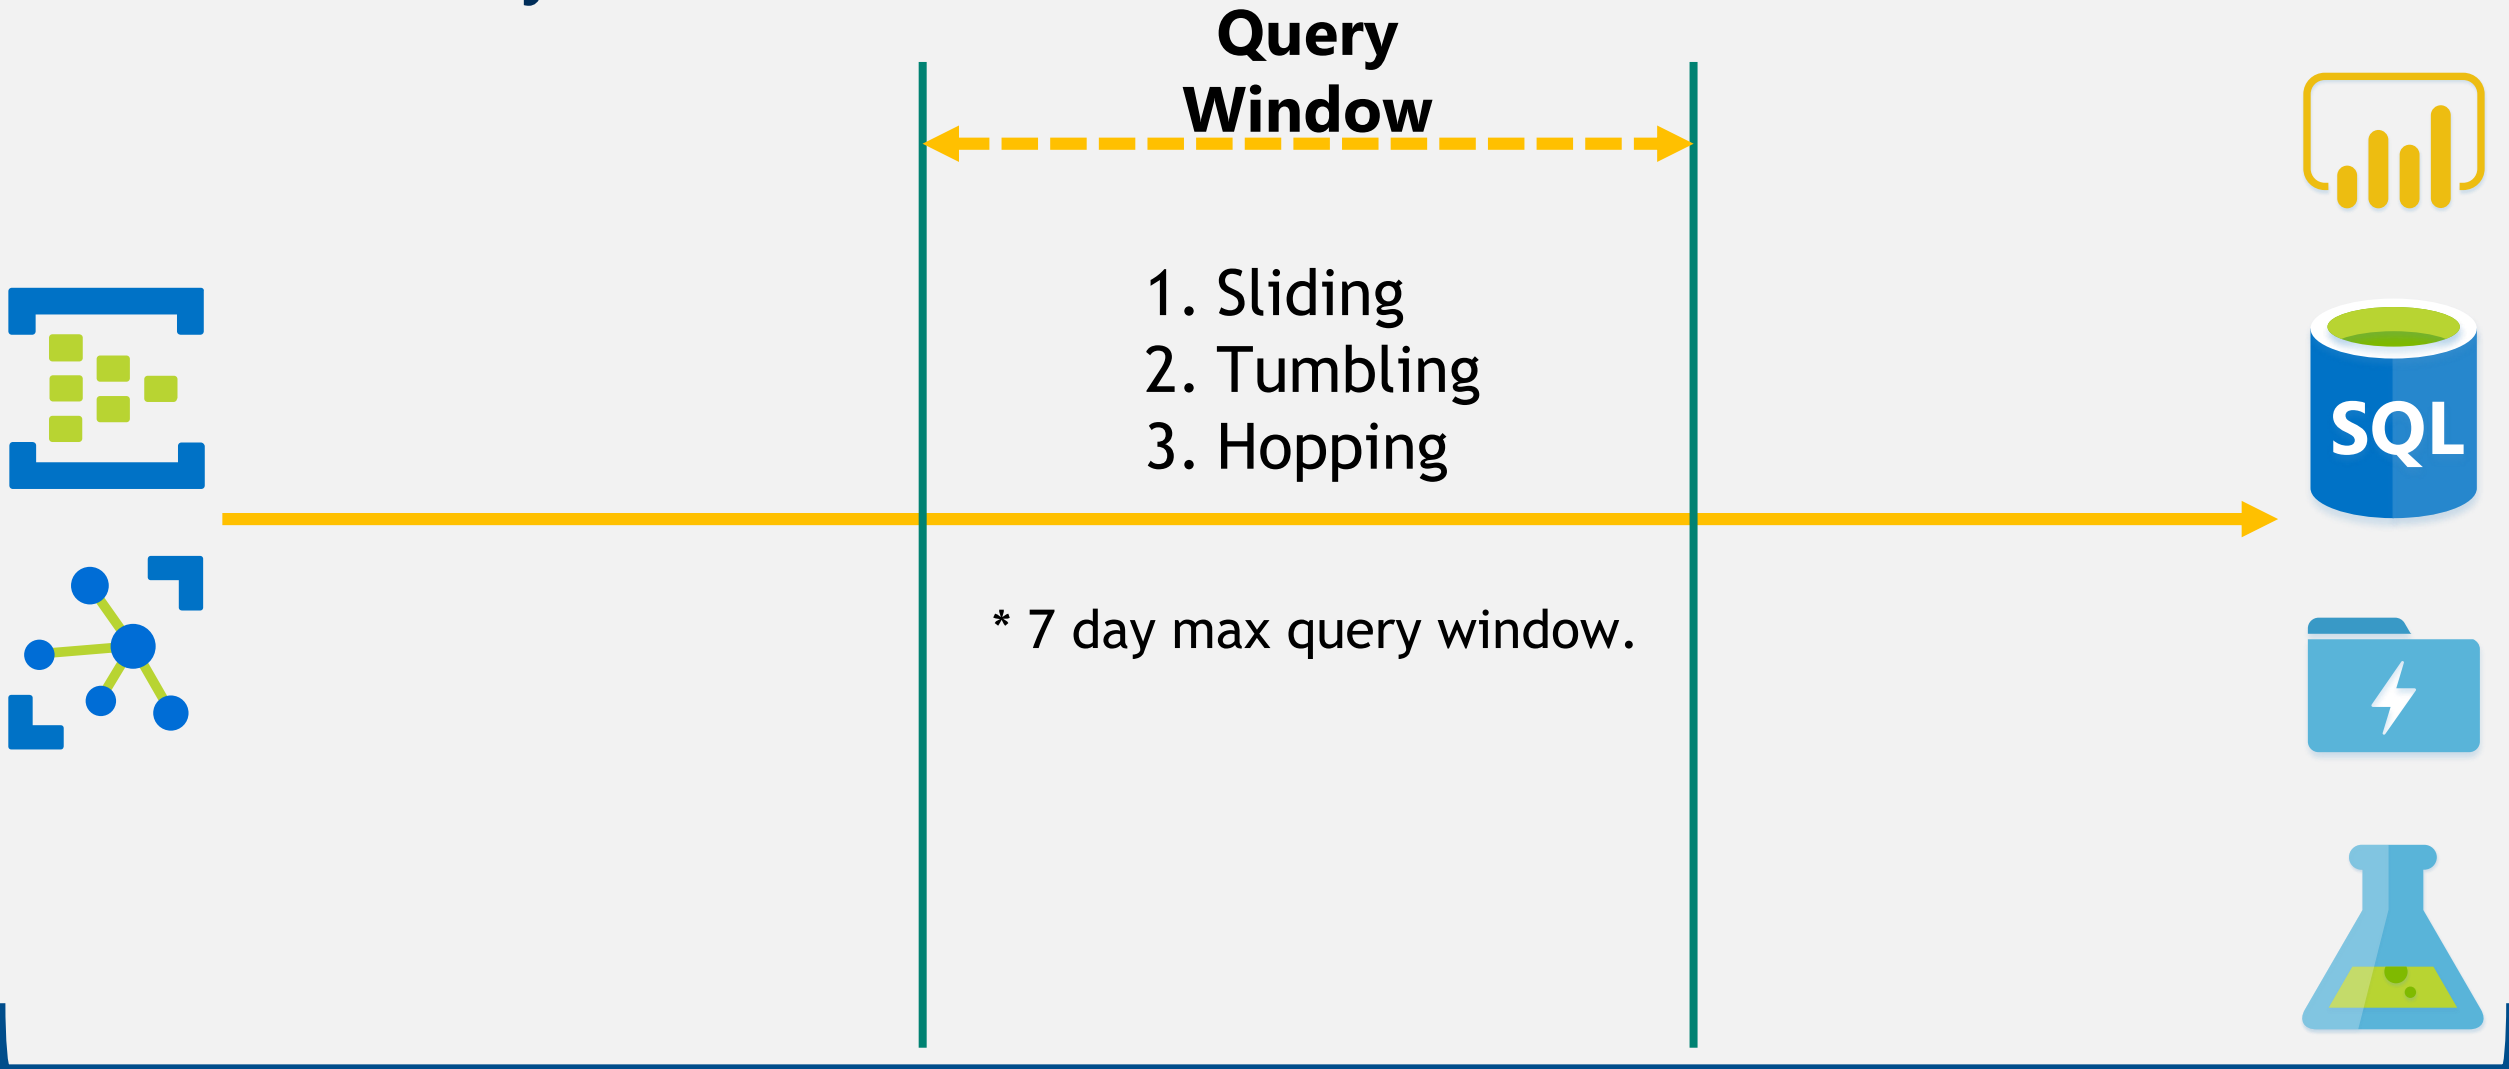


Inputs

Outputs



Azure Stream Analytics

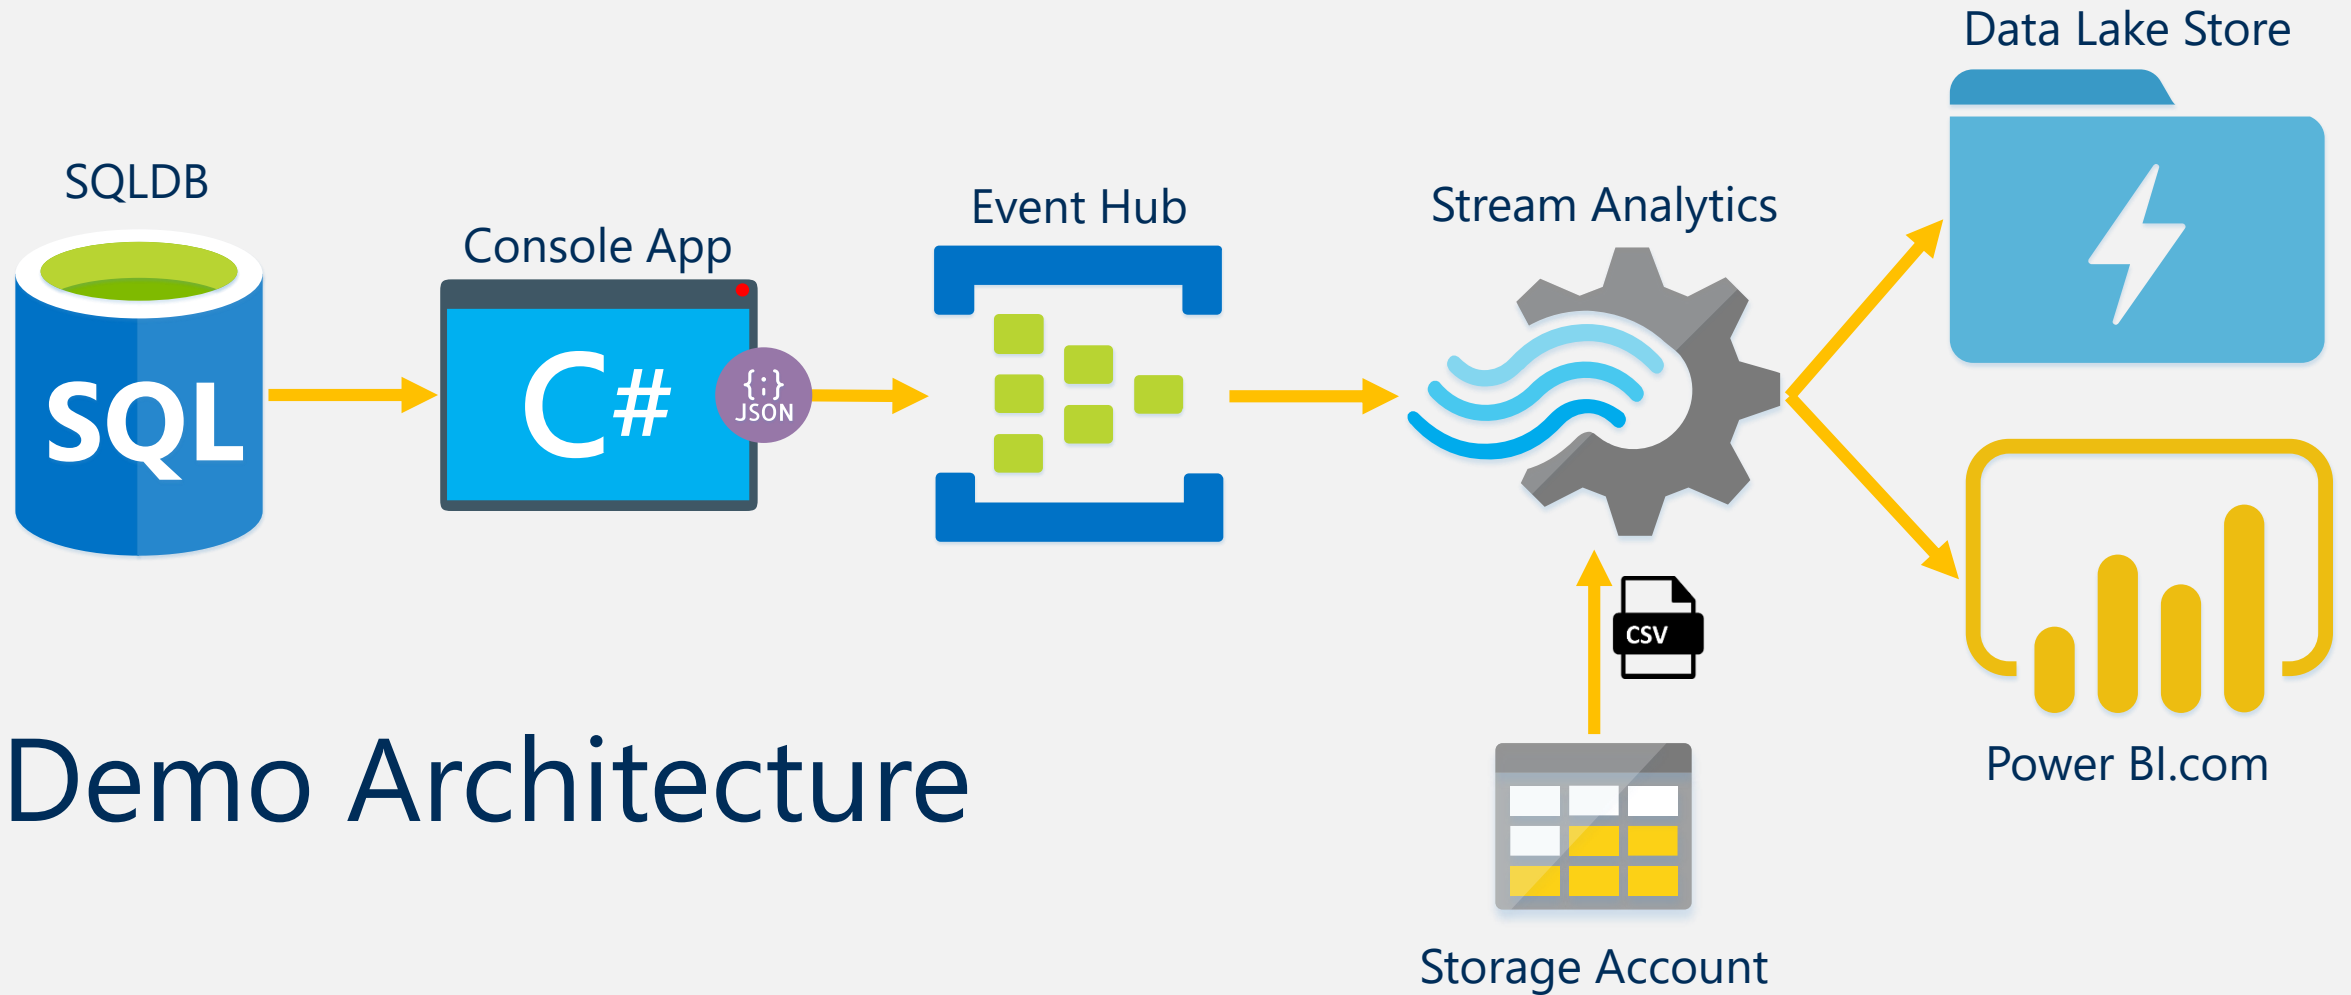


Inputs

<https://msdn.microsoft.com/en-us/library/azure/dn835019.aspx>

Outputs

Demo



Azure Stream Analytics




Real-time data problems

What is *ASA* and why use it

Production Considerations

Lambda Architecture

Production Considerations



Production Considerations



Azure

- 1. Decouple streaming aggregations from persisted storage with multiple jobs/services.



On
Premises



Production Considerations



2. Azure Event Hub Service Endpoints are only accessible via Express Route to on premises resources.



On Premises



VPN



Express Route



VNet



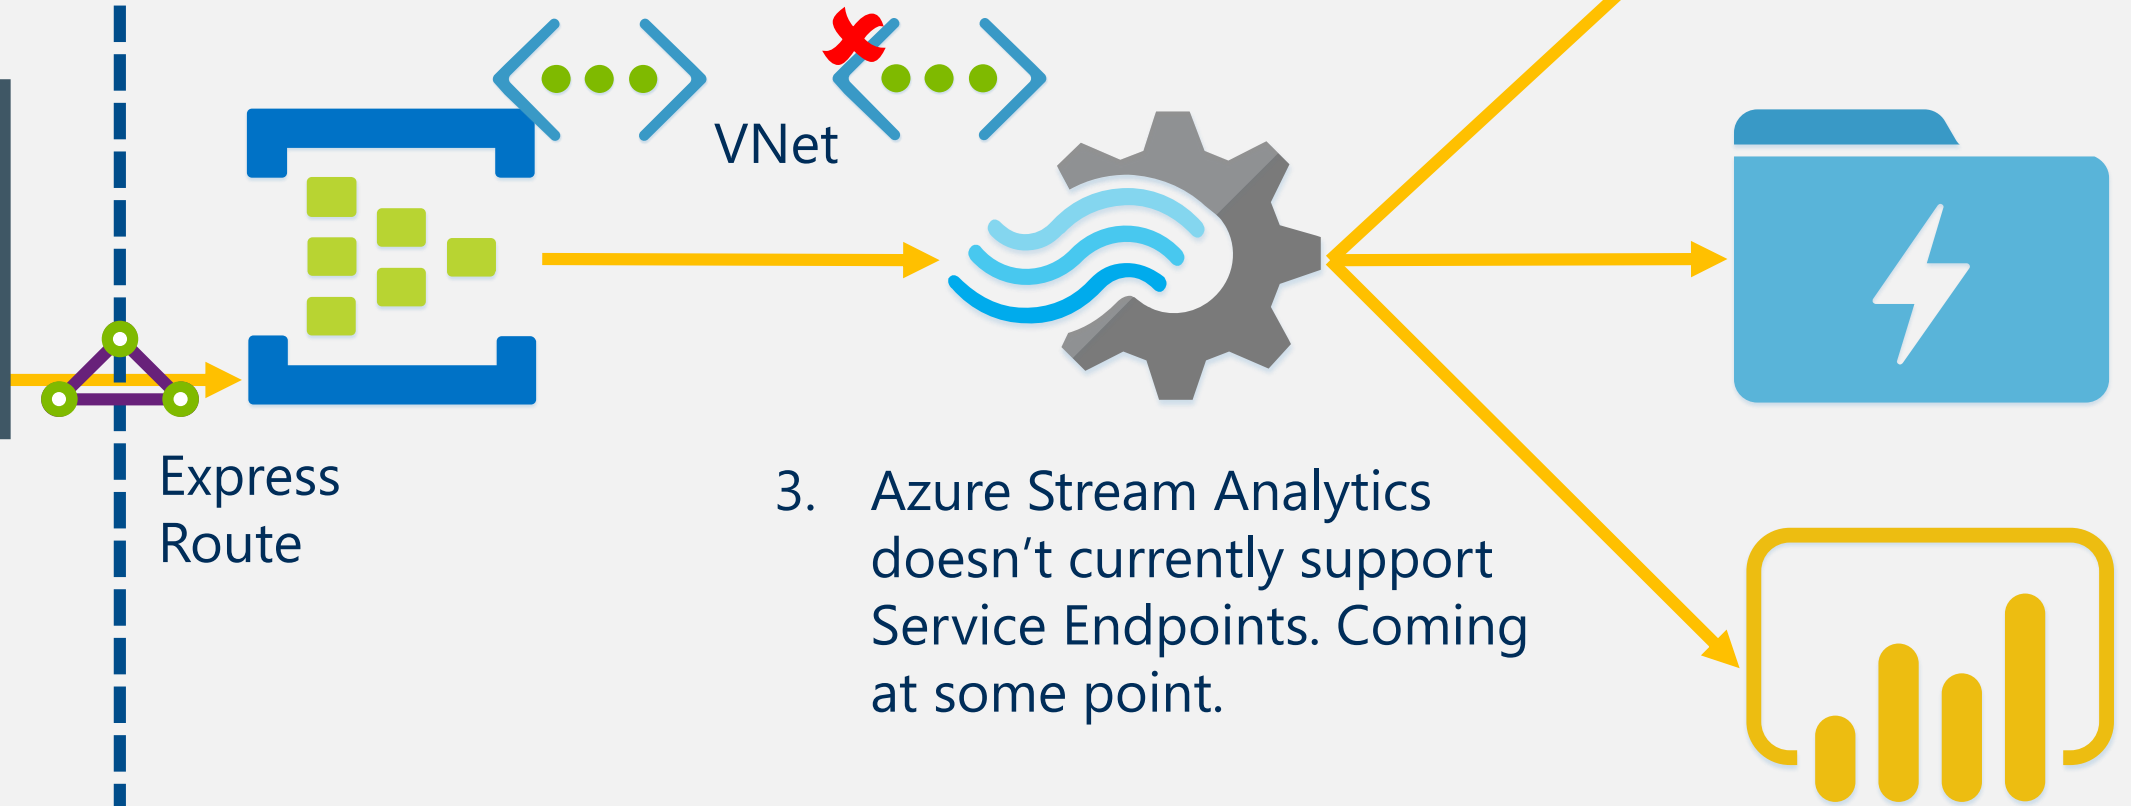
SQL



Production Considerations



-
2. Azure Event Hub Service Endpoints are only accessible via Express Route to on premises resources.



-
-
3. Azure Stream Analytics doesn't currently support Service Endpoints. Coming at some point.

Production Considerations



2. Azure Event Hub Service Endpoints are only accessible via Express Route to on premises resources.



On Premises

Express Route



VNet



VNet

3. Azure Stream Analytics doesn't currently support Service Endpoints. Coming at some point.

4. What other services in your solution need to use Service Endpoints?

Production Considerations




On
Premises





Capture



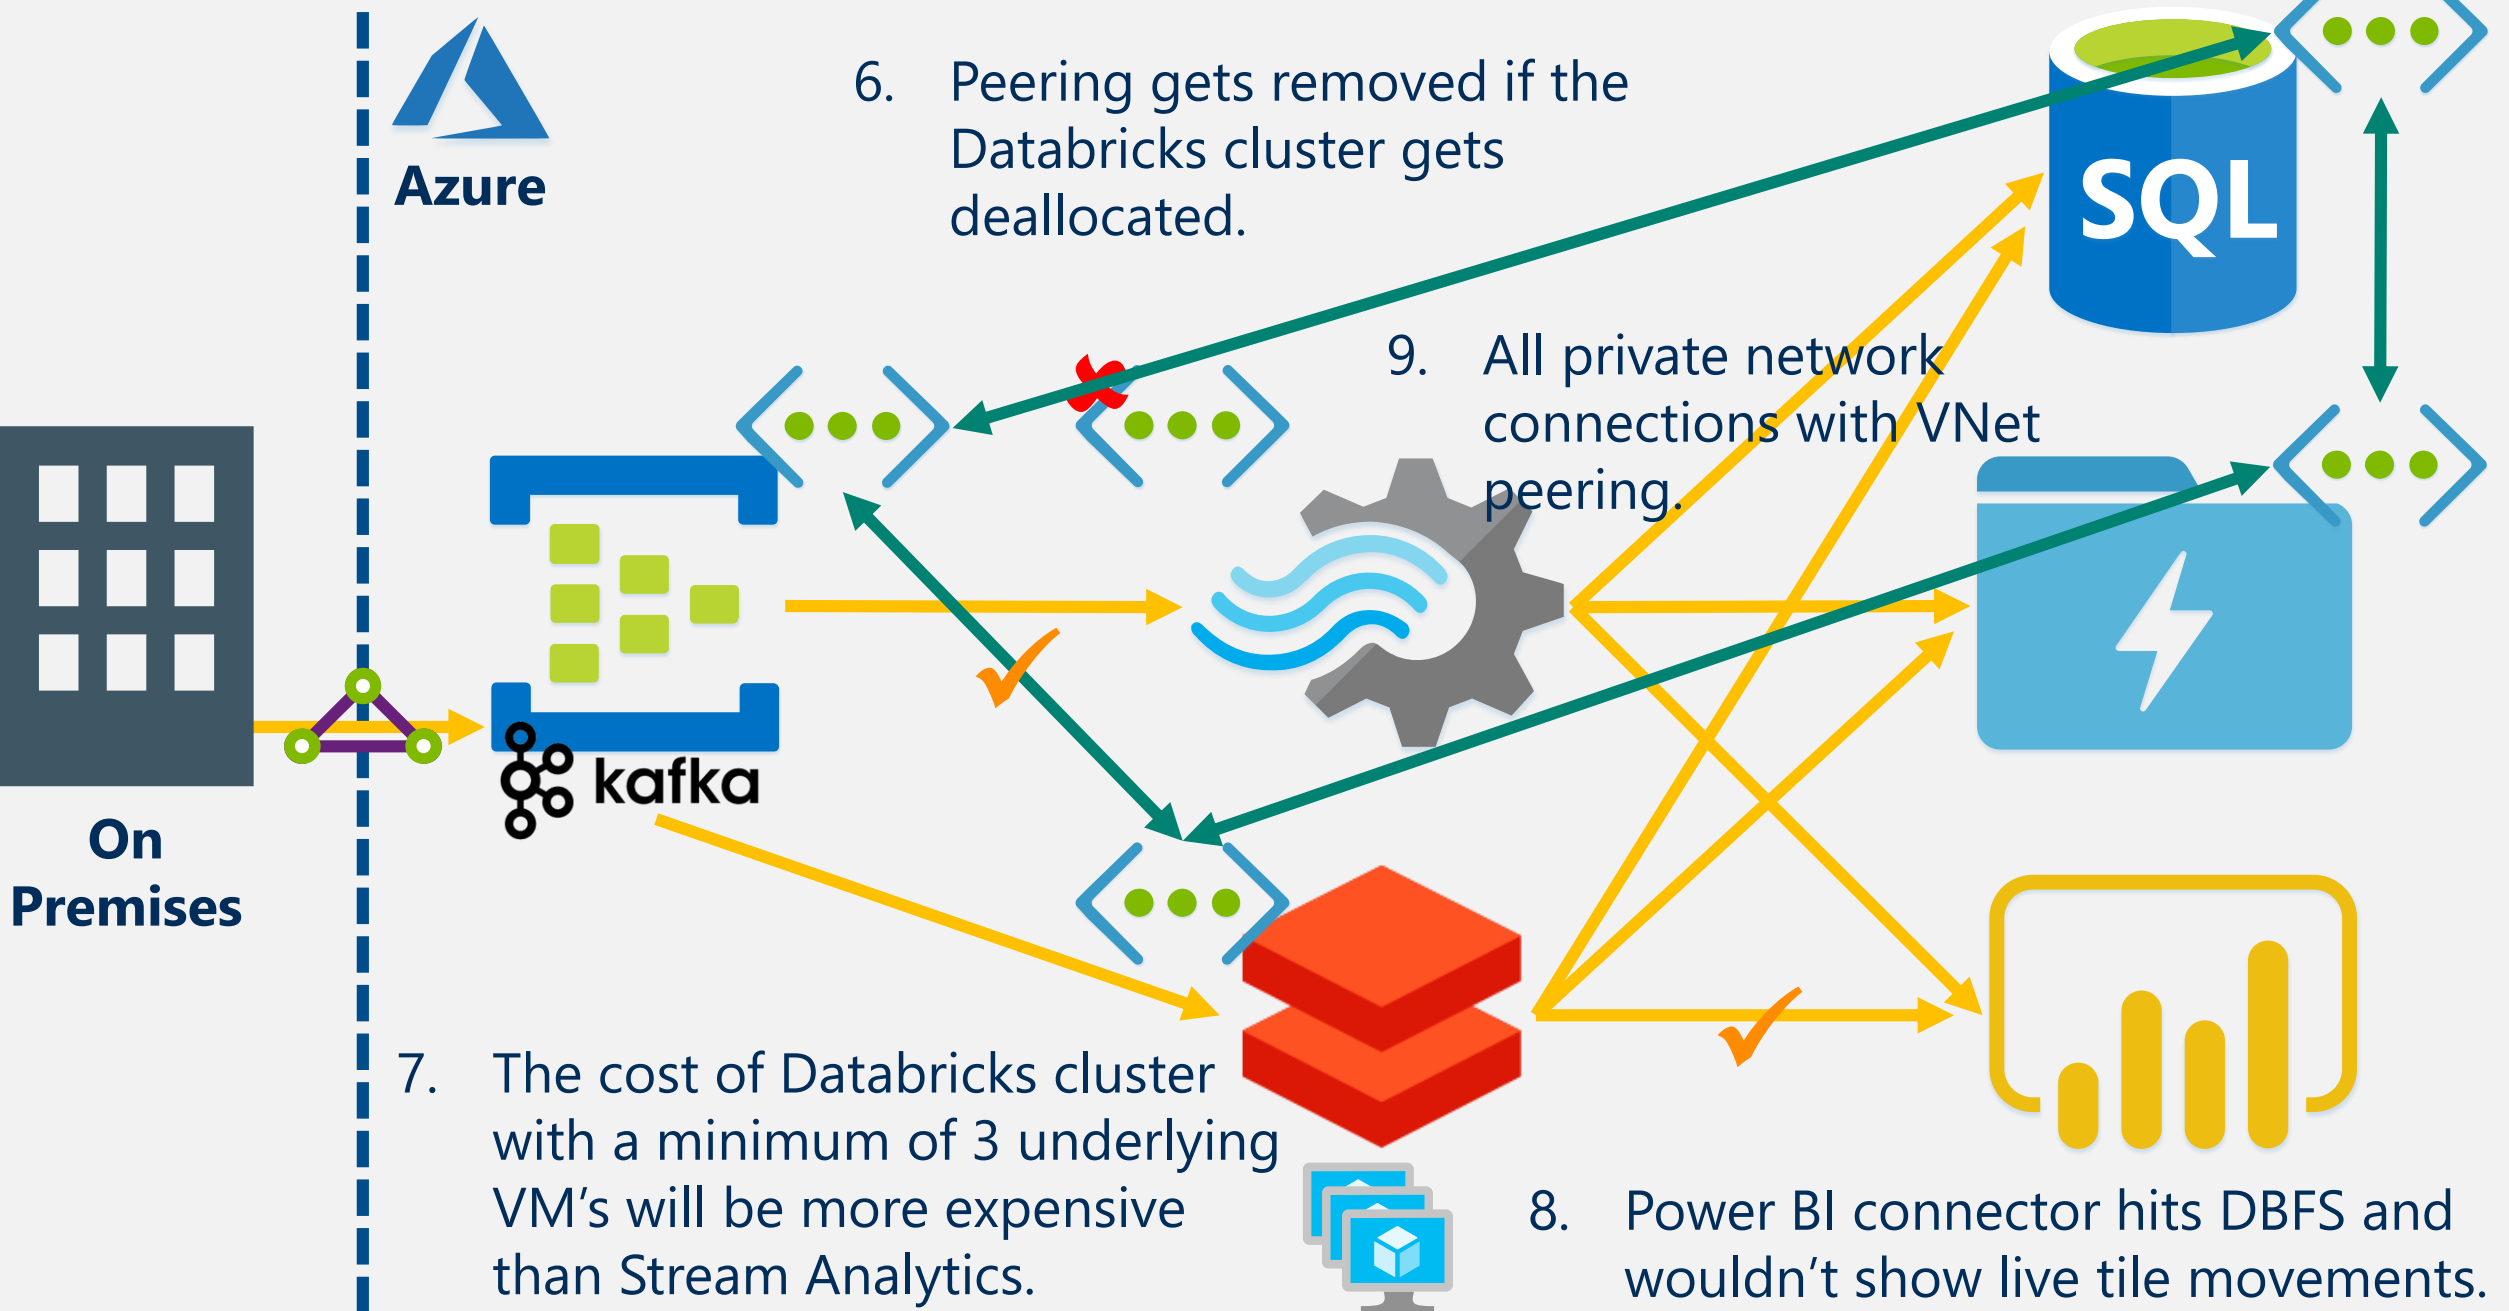
VNet



- 5. Limited flexibility with the target folder structure and no ability to query streamed data before its persisted to storage using capture.



Production Considerations



Production Considerations Summary

1. Decouple streaming aggregations from persisted storage with multiple jobs/services.
2. Azure Event Hub Service Endpoints are only accessible via Express Route to on premises resources.
3. Azure Stream Analytics doesn't currently support Service Endpoints. Coming at some point.
4. What other services in your solution need to use Service Endpoints?
5. Limited flexibility with the target folder structure and no ability to query streamed data before its persisted to storage using capture.
6. Peering gets removed if the Databricks cluster gets deallocated.
7. The cost of Databricks cluster with a minimum of 3 underlying VM's will be more expensive than Stream Analytics.
8. Power BI connector hits DBFS and wouldn't show live tile movements.
9. All private network connections with VNet peering.

<https://docs.microsoft.com/en-us/azure/event-hubs/event-hubs-service-endpoints>

Azure Stream Analytics



Real-time data problems


What is *ASA* and why use it

Production Considerations

Lambda Architecture

Lambda Architecture

Use Batch and Stream technologies together to balance latency, throughput and fault-tolerance




Nathan Marz
& James Warren

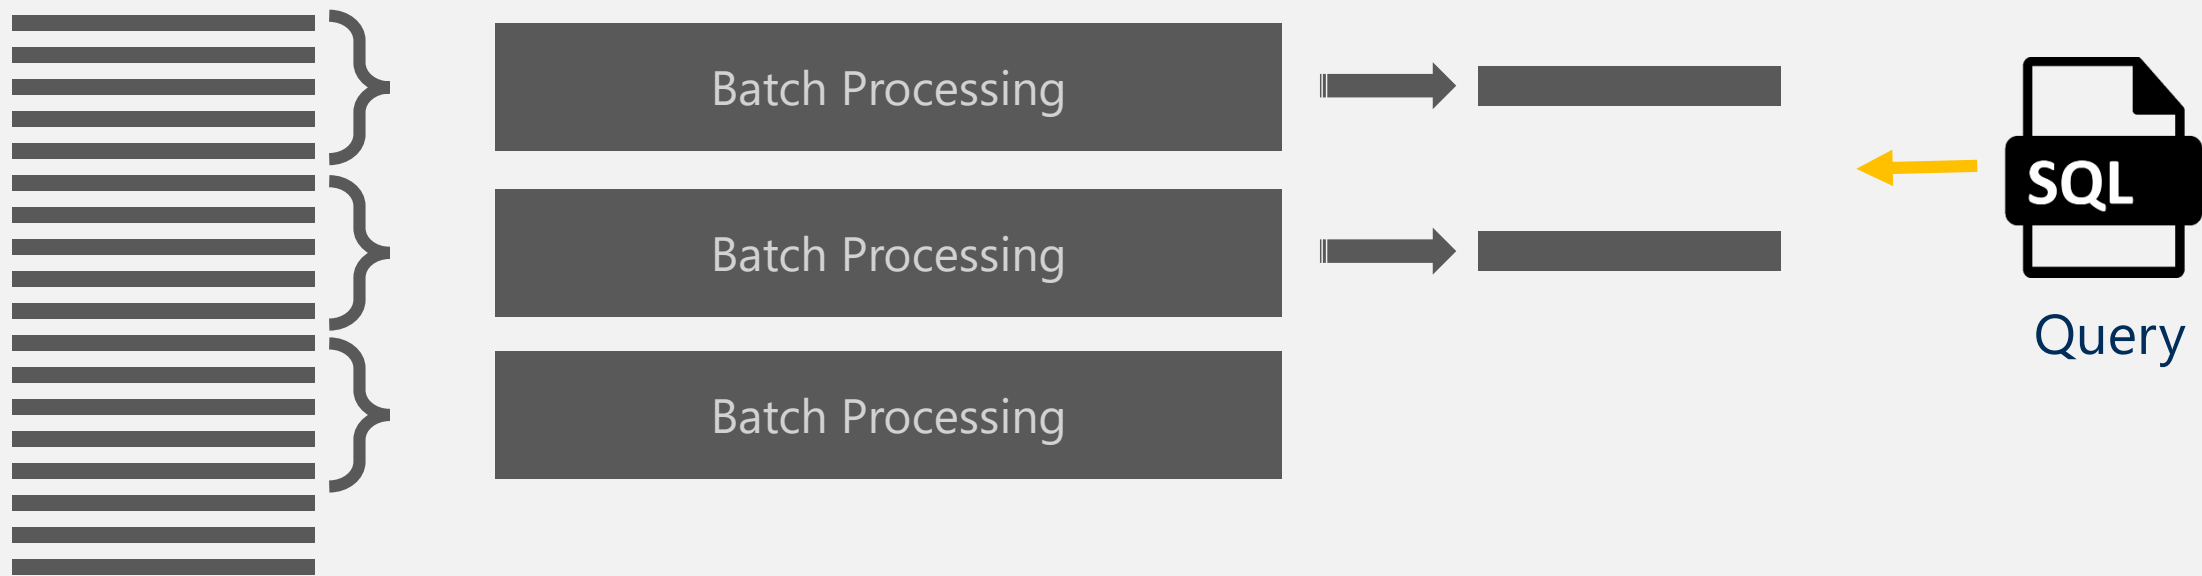


* Pages 14 to 20

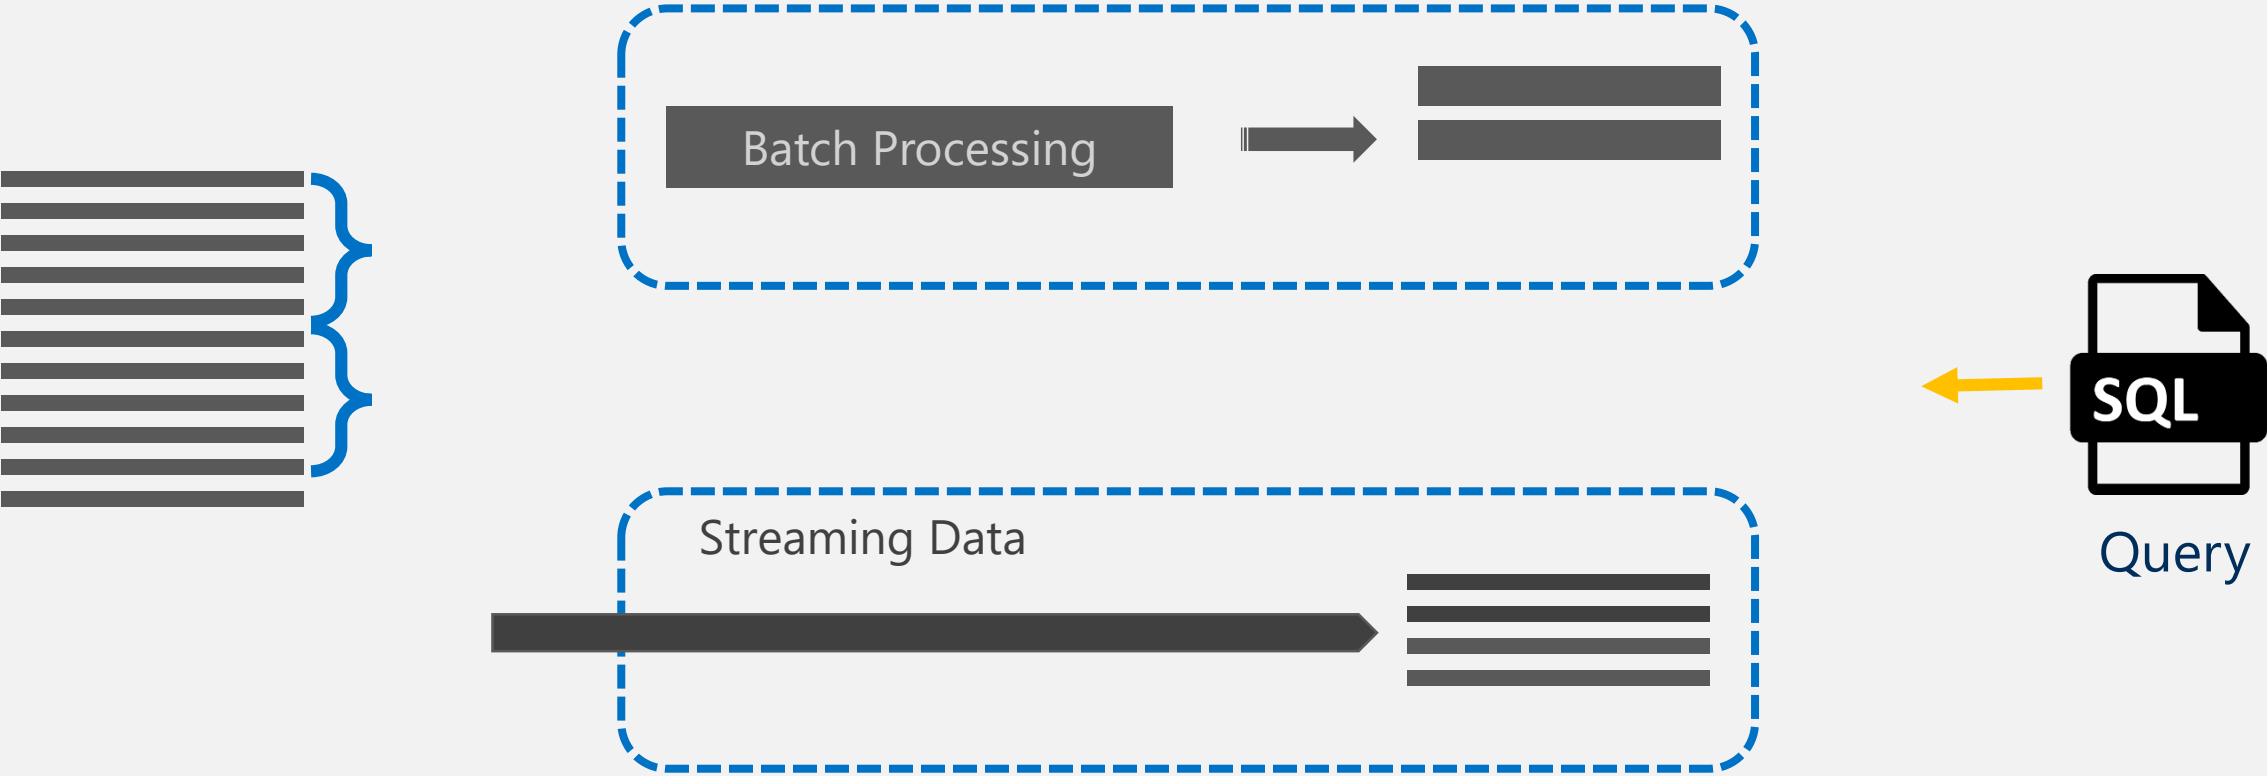
Problem: Timely Data Insights



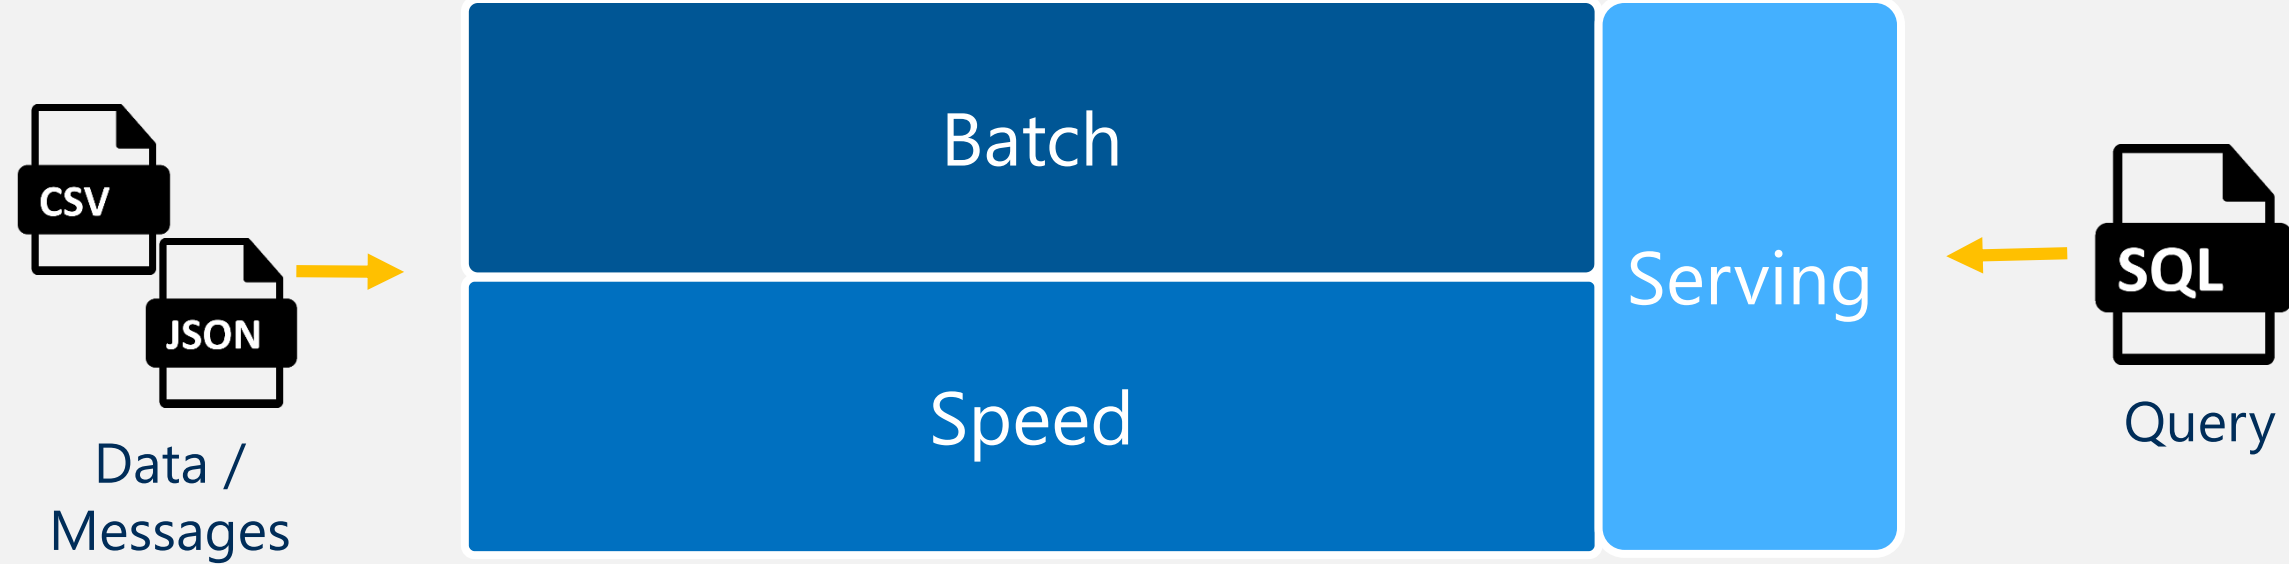
Problem: Timely Data Insights



Solution




Lambda Architecture




The Marz Lambda Architecture




Applying a Lambda Architecture in Azure



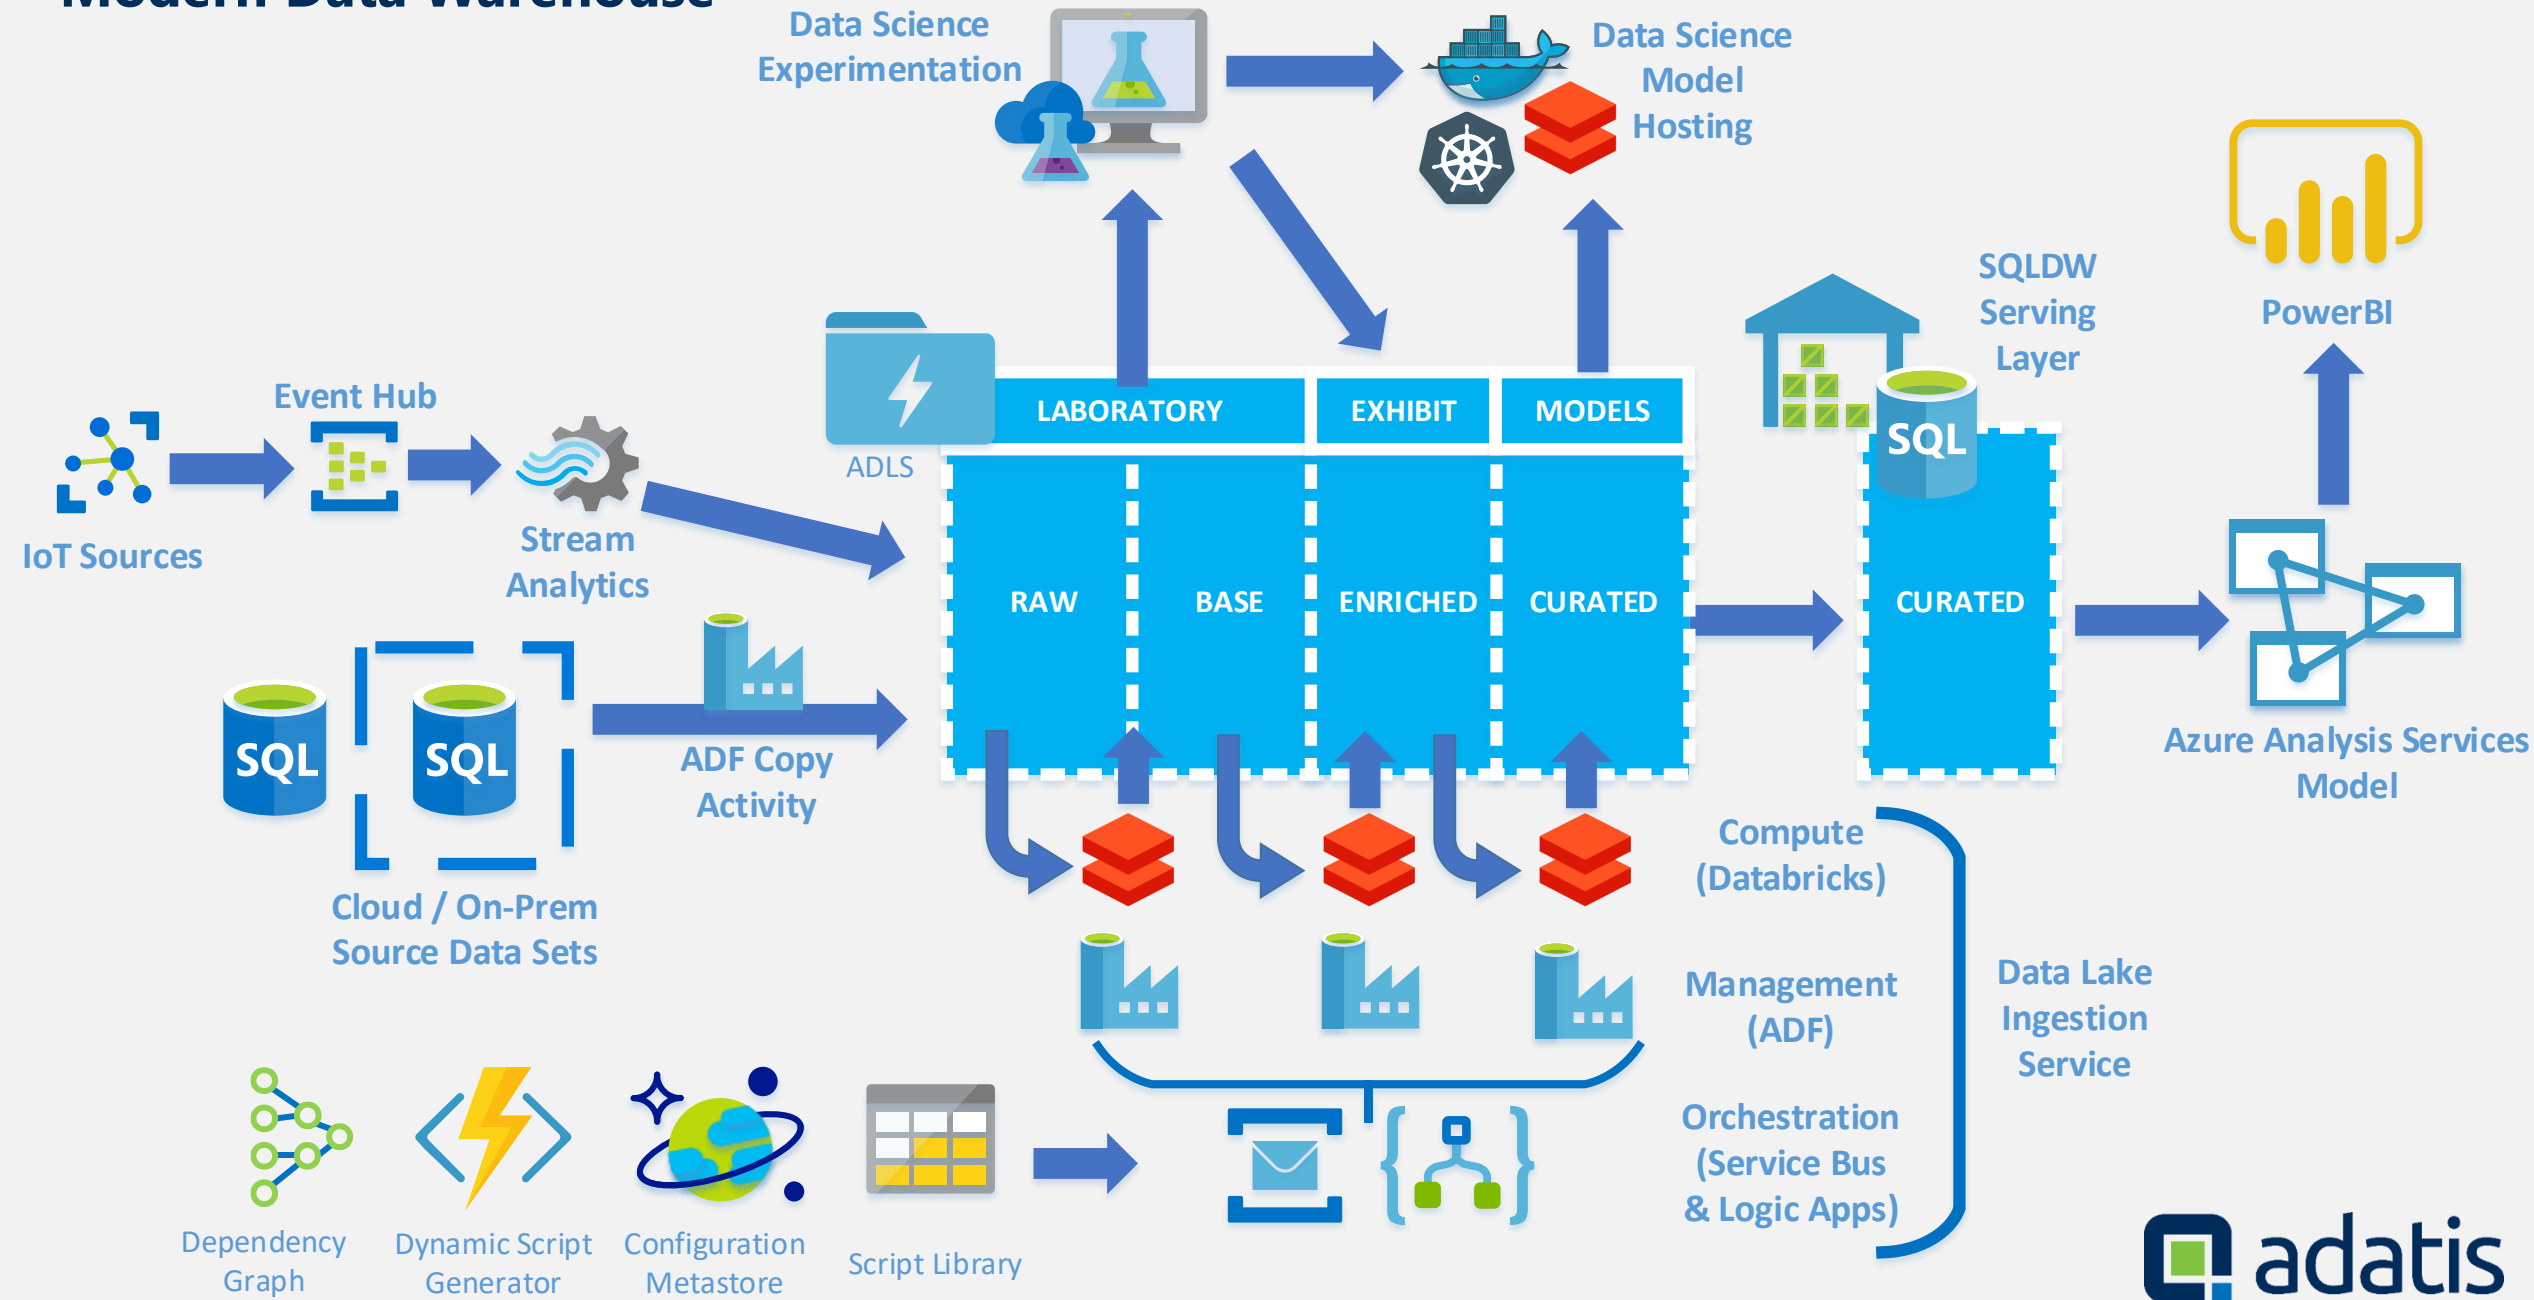
Applying a Lambda Architecture in Azure



Applying a ~~Lambda~~ Architecture in Azure ^ Databricks Delta



Modern Data Warehouse



Azure Stream Analytics



Real-time data problems



What is ASA and why use it



Production Considerations



Lambda Architecture



Thanks for Listening

Paul Andrew

 @MrPaulAndrew



Email: paul@mrpaulandrew.com

Blog: mrpaulandrew.com

GitHub: github.com/mrpaulandrew